

FILE MANAGEMENT

LECTURE 4 / FIT2100 / SEMESTER 2 2019

WEEK 4



FILE MANAGEMENT

INTRODUCTION

LEARNING OUTCOMES

- Understand the basic concepts of **files** and **file systems**
- Understand how the OS boots from a hard drive
- Discuss approaches used in real-world file systems
- Describe the low-level Unix system calls used to **access** and **manipulate** the filesystem

READING

- Stallings: Chapter 12
- Extra information about Windows FAT filesystem:
<https://social.technet.microsoft.com/wiki/contents/articles/6771.the-fat-file-system.aspx>

WHAT IS A FILE?

- A file is a sequential list of characters (bytes)
- Interpretation of data is determined by the application program
- Text file:
 - File is made up of characters
 - Each byte is assigned a character code value
 - When printing the file contents, the terminal displays actual characters
- Binary files
 - Bytes are used to store values directly
 - e.g. 4 bytes can be used to represent a 32 bit integer
 - 1 byte can store a value from 0 to 255.
 - Byte values do not *need* to match readable character codes
 - Trying to view the file as *text* shows nonsense characters

WHAT IS A FILESYSTEM?

SOFTWARE WITHIN THE OPERATING SYSTEM

- Manages data on the disk
 - Translates block data on the disk into abstract concepts of files and directories
- Provides services for accessing and manipulating files

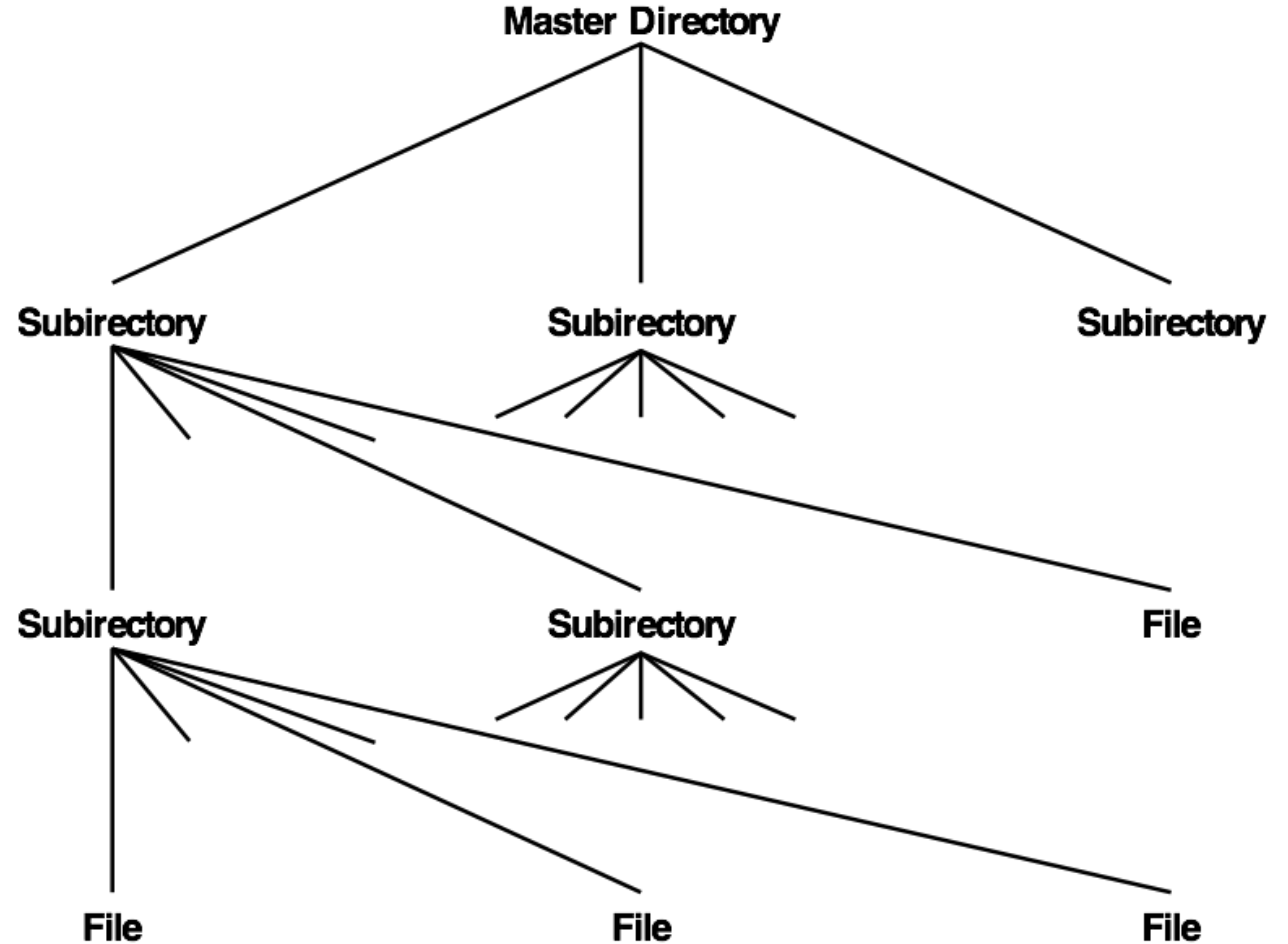
STRUCTURED DATA STORED ON THE DISK ITSELF

- Multiple approaches exist
- Different operating systems use different filesystems
- Two most common approaches
 - File Allocation Tables (Windows)
 - Inode-based systems (UNIX).

HIERARCHICAL FILESYSTEMS

AKA 'TREE-STRUCTURED DIRECTORY'

- Multiple file-organisation approaches are possible
- Tree-structured or 'hierarchical' is most common
- A 'root directory' or 'master directory' contains files and sub-directories.



REAL-WORLD FILE SYSTEMS (FAT)

The IBM Personal Computer DOS
Version 1.00 (C)Copyright IBM Corp 1981

*.com

COM	1920	07-23-81
COM	6400	08-13-81
COM	3231	08-04-81
COM	2560	08-04-81
COM	1395	08-04-81
	896	08-04-81
	1216	08-04-81
	1124	08-04-81
	1620	08-04-81
	252	08-04-81
	250	08-04-81
	860	08-04-81
	3392	08-04-81
	49	08-04-81
	0	08-04-81
		08-04-81

FAT

FILE ALLOCATION TABLE

COMMON ON WINDOWS SYSTEMS (FAT32, ETC.)

- Reference: <https://social.technet.microsoft.com/wiki/contents/articles/6771.the-fat-file-system.aspx>
- Drive is split into equal-size **blocks** (FAT32: a 1TB hard drive is split into 32KB 'clusters')
- A file may occupy one or more blocks
- A **table** at the start of the disk contains an entry for every block on the drive.
- In case the table is corrupted, files can no longer be located.
 - Two copies of the table are kept, for redundancy


Boot sector	File allocation table 1	File allocation table 2 (duplicate)	Root directory	Other directories and all files
-------------	-------------------------	-------------------------------------	----------------	---------------------------------

FRAGMENTATION

TYPES OF FRAGMENTATION


EXTERNAL FRAGMENTATION

- External fragmentation
- Free blocks between files, but the spaces are not large enough to fit certain files
 - e.g. a small file was deleted, leaving a space which is too small to fit a new larger file.
- **Fragments** of wasted space unless the file can be split over multiple fragments.
- Was a problem in very old filesystems.



INTERNAL FRAGMENTATION


- If a file is smaller than a single block
- Unused space inside the block
 - Cannot be allocated to other files



HOW A DIRECTORY WORKS (FAT)

DIRECTORY = INDEX OF FILE LOCATIONS ON DISK

- A directory stores the address of the **starting block** of each file.
- Files are allocated to the first empty block.
- If the file is too large for one cluster, available clusters are chained together
 - Each block stores a reference to the next one where the file continues
 - By following the chain, the entire file can be read



FILE SYSTEM AGEING

FAT

NO EXTERNAL FRAGMENTATION

- Chaining blocks together means that no blocks are wasted
 - File does not need to be contiguous on the disk
- Files are added and deleted, creating gaps of free sectors
- A large file may not fit into a single gap
 - File is spread over different tracks
- The hard drive head must move to different locations to access the entire file
 - Access times get slower over time
- Disk requires periodic '**defragmentation**' to improve performance.
 - Files are rearranged into contiguous blocks

Defragmenting disk

Creative commons attribution:
XZise, Wikimedia Commons

1		3	1	2	3	5	1		1
2	3	8	4			3	4		
2		5	4		4	6			6
	5		2	4		3		2	
9	7		6	4	6		7	1	
5		3		2		1			

REAL-WORLD FILE SYSTEMS (UNIX)

TYPES OF FILES IN UNIX

Regular, or ordinary

- contains arbitrary data in zero or more data blocks

Directory

- contains a list of file names plus pointers to associated inodes

Special

- contains no data but provides a mechanism to map physical devices to file names

Named pipes

- an interprocess communications facility

Links

- an alternative file name for an existing file

Symbolic links

- a data file that contains the name of the file it is linked to


UNIX FILESYSTEMS: INODES

- All types of UNIX files are stored in a table of **inodes**
 - ‘Index nodes’
 - An alternative system to file allocation table (FAT) systems
- The inode for a file stores key information about that file
- Each file is controlled by exactly one inode
- Several filenames may be associated with a single inode
 - The same physical file can be mapped to multiple directories.

HOW DO DIRECTORIES WORK?

INODE-BASED FILESYSTEMS

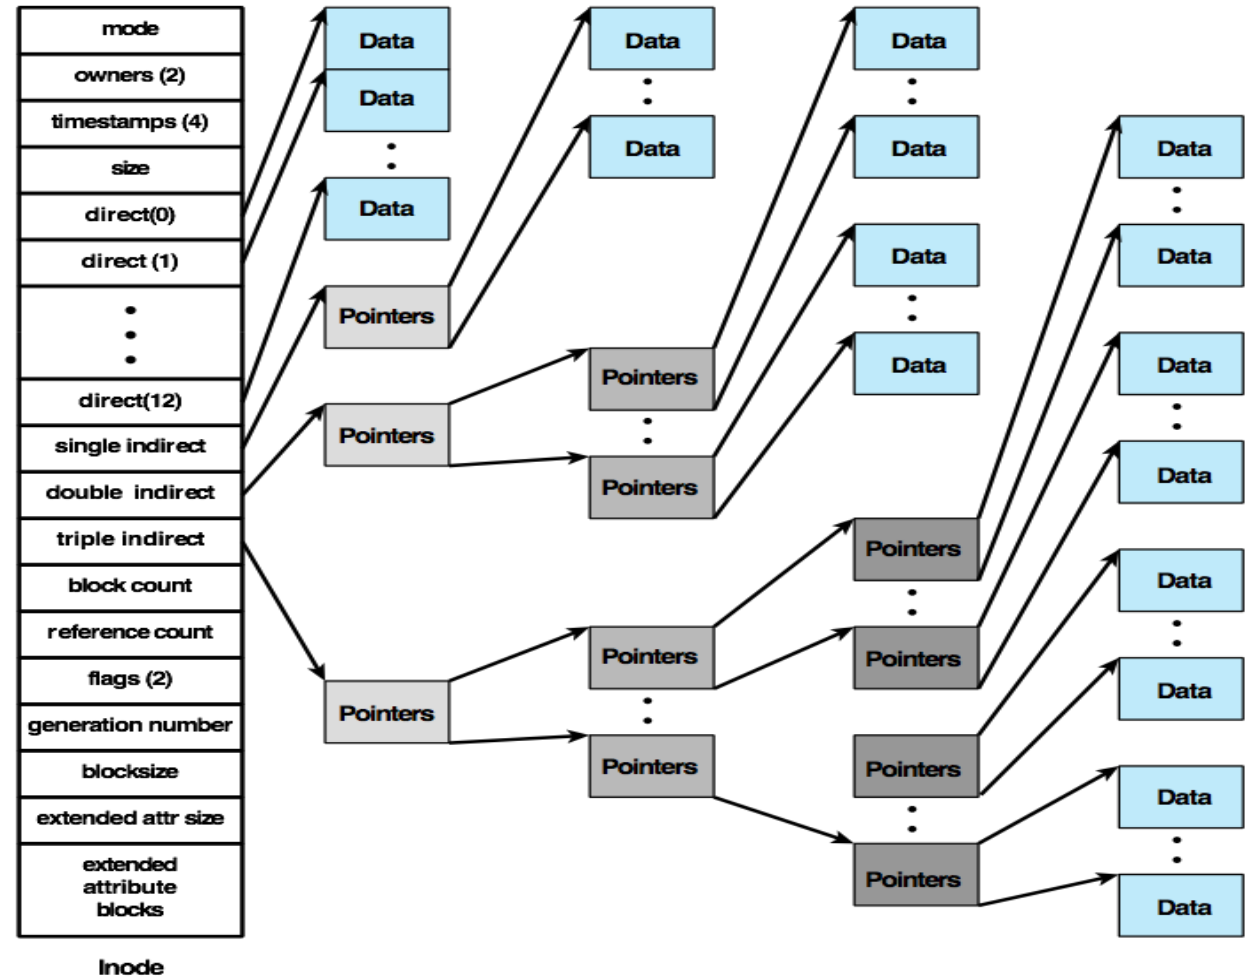
- A directory is stored like a regular file
- Directory contains a table of inode numbers and filenames
- Files are located by looking up the appropriate inode in the inode table.



UNIX: FILE SYSTEM STRUCTURE

INODE AND FILE

- Each inode contains various information about a file, and how to locate it on disk.
- When a file is opened, its corresponding inode is loaded into main memory.
- If a file requires 12 blocks on disk or fewer, these addresses are listed **directly** in the inode itself.
- If the file is larger than 12 blocks, the filesystem allocates special **pointer blocks** containing the remaining addresses
- If the file is very large, multiple levels of pointer blocks are used



BOOTING THE OPERATING SYSTEM



BOOTING THE O.S.

AKA 'BOOTSTRAPPING'

A PROBLEM

- Before the computer is turned on, the OS lives only on the hard drive
- But, without the OS running, how do you locate files on the hard drive?

BOOTSTRAPPING

- The idea of using a small, primitive system to start up a slightly more advanced system.
- Many approaches exist, we will look at a common one on PCs.

BIOS

BASIC INPUT OUTPUT SYSTEM

PRIMITIVE SOFTWARE ON A HARDWARE CHIP

- EEPROM
 - Electrically **E**raseable **R**ead **O**nly **M**emory
- A tiny area at the very end of main memory that is **non-volatile**
- Machine code is already here when the power is turned on
 - CPU always looks here for the first code to execute
- Contains instructions for recognising connected hard drives
- Can remember hard drive configuration using a small battery inside the computer
- Does not need to know how to read a filesystem, only how to load raw bytes from the hard drive into memory



**EEPROM module
from a typical PC**

Creative Commons
attribution: Nevit
Dilmen

MASTER BOOT RECORD

MBR – ALSO KNOWN AS BOOT SECTOR


TRACK 0, SECTOR 0

- The first 512 bytes on a bootable hard drive
- Contains machine code, known as 'bootstrap program'
- The computer's BIOS loads this into main memory and points the CPU to it – transfers control.
- **Contains instructions for locating and loading the kernel, which loads the rest.**

EXISTS INDEPENDENTLY OF THE FILESYSTEM

- The MBR is not a 'file'
 - Just raw bytes read directly from start of the hard drive
- If this OS hasn't been loaded yet, there are no services available for locating files on the rest of the drive.

TRACK 0
SECTOR 0



ACCESSING THE FILESYSTEM IN LINUX

SYSTEM CALLS FOR FILE I/O

OS PROVIDES ACCESS TO THE FILESYSTEM VIA SYSTEM CALLS

- **open(...)**
 - The operating system allocates a number (file descriptor) to each open file within a running process
 - The OS also stores the current read/write position within the file
- **write(...)**
 - The program tells the OS the starting address in main memory and how many bytes to write into the file
- **read(...)**
 - The program gives the OS a main memory address to read file contents into, and how many bytes may be read into memory (i.e. array size).
 - The operating system increments the file position after each read/write.

SYSTEM CALLS (2)

- **lseek(...)**
 - The OS changes the current read/write position to a different place within the file
- **close(...)**
 - The OS de-allocates the file descriptor information from memory.
- Many other system calls include: **creat()**, **link()**, **unlink()**, **chmod()**, **chdir()**, **lchown()**, **stat()**, **mkdir()**, etc.

SUMMARY (LECTURE 4)

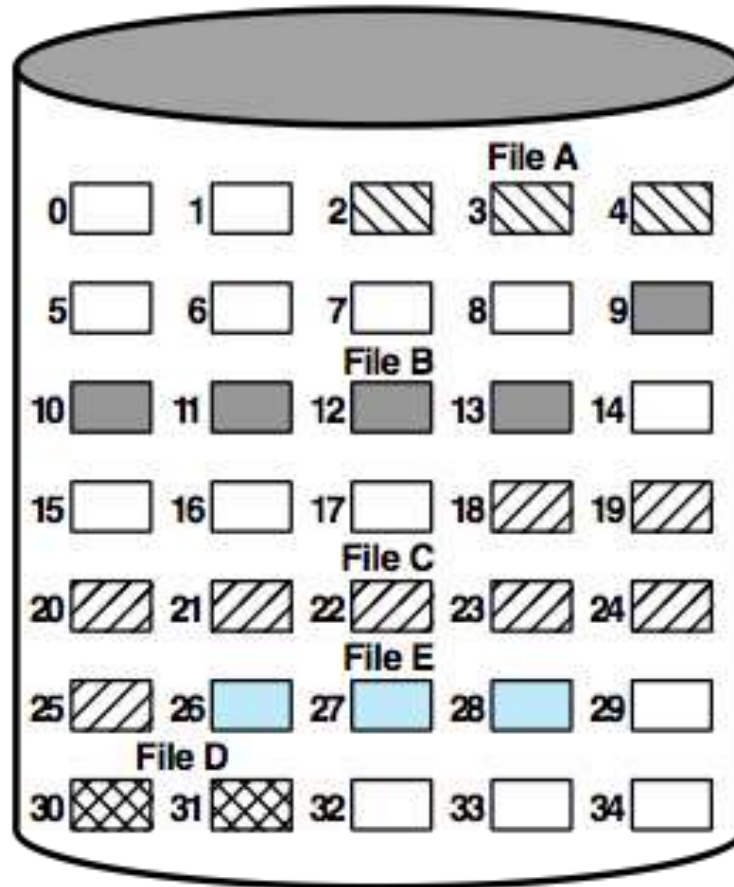
FILE MANAGEMENT

- A file management system is a set of system software that provides services to users, and structured data on the drive itself
- Various approaches to file management exist
 - File allocation table systems
 - Inode-based systems
- The computer must be able to boot the operating system from the hard drive even before the filesystem software has been loaded.
- **Next week: processes**

SUPPLEMENTARY SLIDES

File Allocation: Contiguous

External
fragmentation

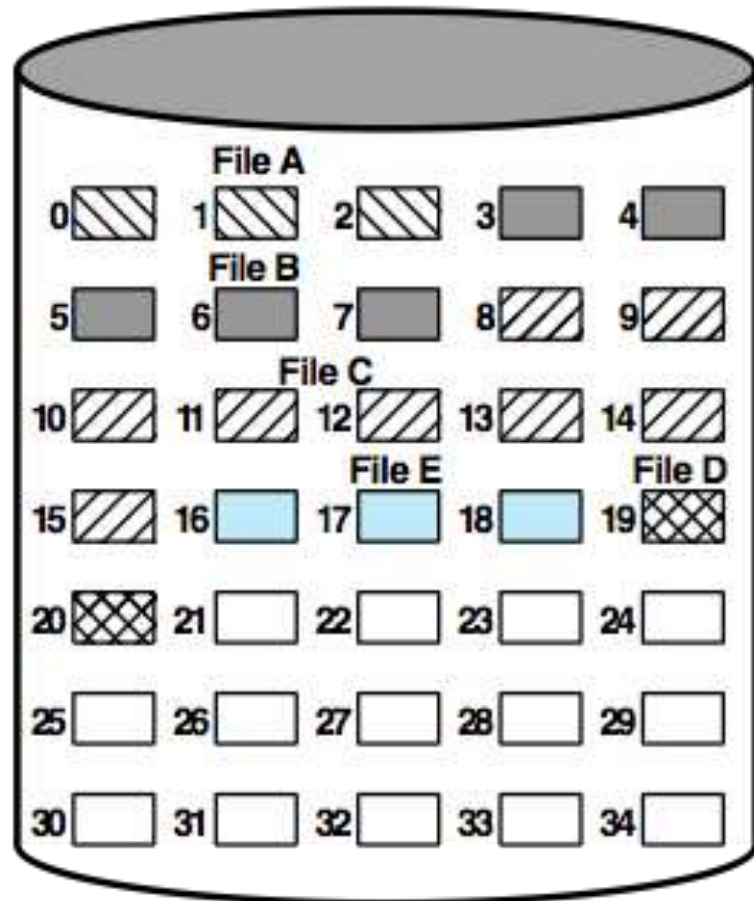


File Allocation Table

File Name	Start Block	Length
File A	2	3
File B	9	5
File C	18	8
File D	30	2
File E	26	3

- A single contiguous block is allocated to a file at the time of file creation
- FAT needs just one single entry for each file, showing the starting block and the length of the file

File Allocation: Contiguous



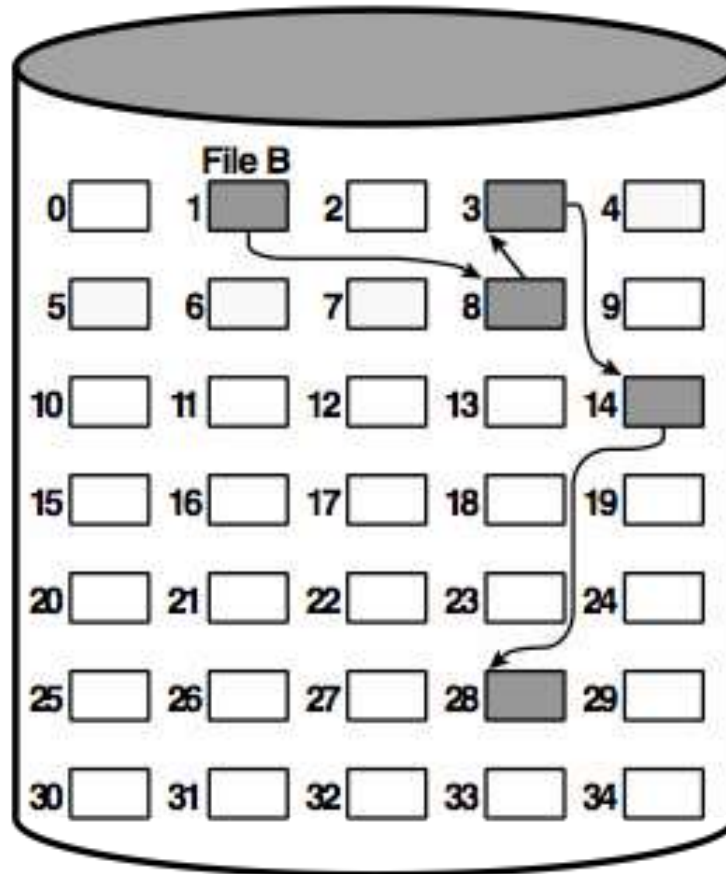
File Allocation Table

File Name	Start Block	Length
File A	0	3
File B	3	5
File C	8	8
File D	19	2
File E	16	3

(after compaction)

File Allocation: Chained

No external fragmentation



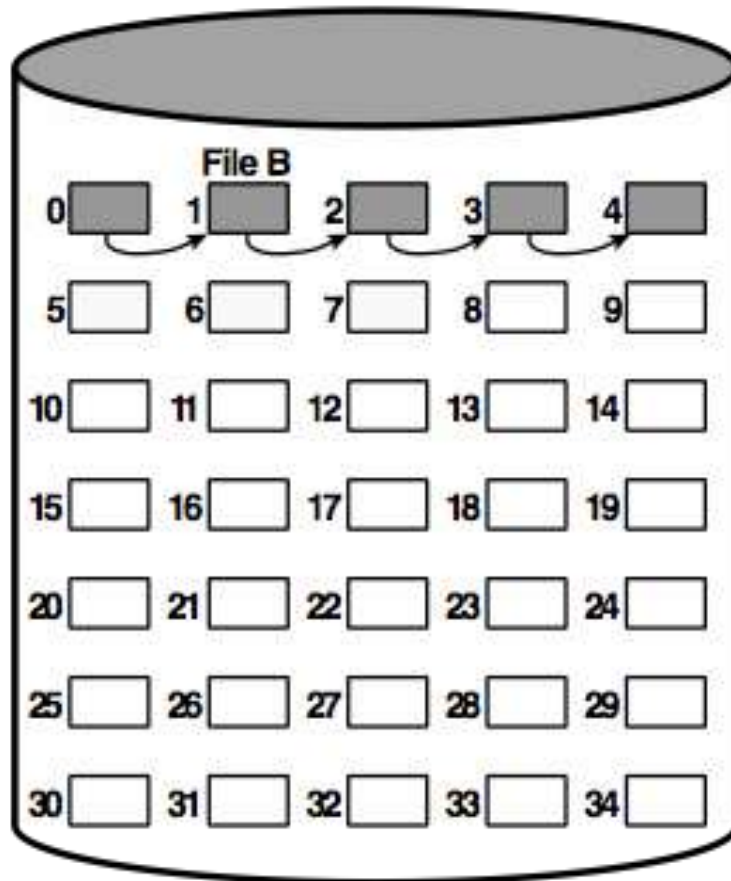
File Allocation Table

File Name	Start Block	Length
...
File B	1	5
...

- Allocation is on an individual block basis
- Each block contains a pointer to the next block in the chain
- FAT needs just a single entry for each file

File Allocation: Chained

No accommodation
for the principle
of locality



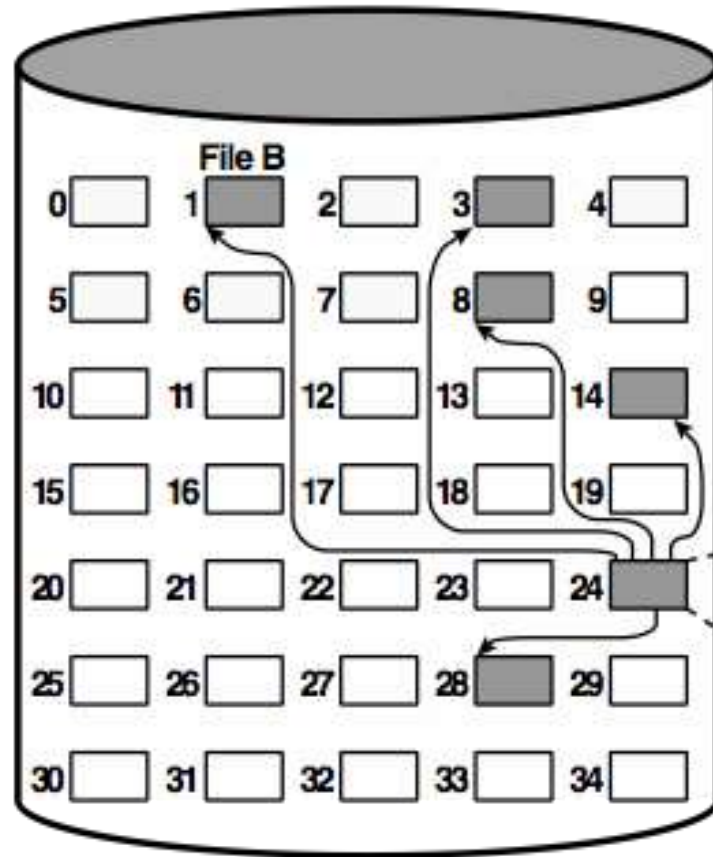
File Allocation Table

File Name	Start Block	Length
...
File B	0	5
...

(after consolidation)

File Allocation: Indexed with Block Portions

Eliminates
external
fragmentation



File Allocation Table


File Name	Index Block
...	...
File B	24
...	...

1
8
3
14
28

- FAT contains a separate 1-level index for each file
- The index has one entry for each portion allocated to the file

File Allocation: Indexed with Variable-Length Portions

Improves
locality



File Allocation Table	
File Name	Index Block
...	...
File B	24
...	...

Start Block	Length
1	3
28	4
14	1

Indexed allocation supports
both sequential and direct
access to the file