

FIT2100 Semester 2 2019

Lecture 7 (Part A): Threads

(Reading: Stallings, Chapter 4)

WEEK 8



Lecture 4 (Part 1): Learning Outcomes

- ❑ Upon the completion of this lecture, you should be able to:
 - Understand the **distinction between process and thread**
 - Describe the basic design issues for threads
 - Explain the difference between **user-level threads** and **kernel-level threads**
 - Discuss thread management in Unix/Linux*

*Reading from Stallings, Chapter 4 (4.6)

WHAT DO WE UNDERSTAND ABOUT PROCESSES?

❑ A process 'owns' resources

- Space in main memory, open files, I/O devices, etc.
- An independent **process image**
- One process is prevented from interfering with another process's resources and image as allocated by the OS

❑ Scheduling and execution

- A process follows a 'path' of execution
- Execution may be interleaved with other processes
- Scheduled and dispatched by the OS.

What is a *thread* for a process?

The Concept of Threads

❑ The unit of dispatching for execution is referred to as:

- Thread or
- Lightweight process

❑ The unit of resource ownership is referred to as:

- Process or task



e.g. Windows or
Unix

The entity that owns a
resource is a process

The Concept of Threads

❑ The unit of dispatching for execution is referred to as:

- Thread or
- Lightweight process

❑ The unit of resource ownership is referred to as:


- Process or task

❑ Multi-threading:

- The ability of an OS to support multiple concurrent paths of execution within a single process
- **1 process : multiple threads of execution.**


Single-Threaded Approaches

- ❑ Left side of figure:
 - A single thread of execution per process.
- ❑ The concept of a thread is not recognised — referred as a **single-threaded** approach.
- ❑ Example: MS-DOS, Windows 3.1



Multi-Threaded Approaches

- ❑ The right half of the figure depicts the **multi-threaded** approach.
- ❑ One process with multiple threads.
- ❑ Example: Windows, Linux.



The Concept of Processes (revisit)

- ❑ The unit of **resource allocation** and a unit of **protection**.
- ❑ A (virtual) address space that holds the process image.
- ❑ Protected access to:
 - ✦ **Processor(s)**
 - ✦ **Other processes**
 - ✦ **Files**
 - ✦ **I/O resources**



**Interprocess
communication**

Threads within a Process

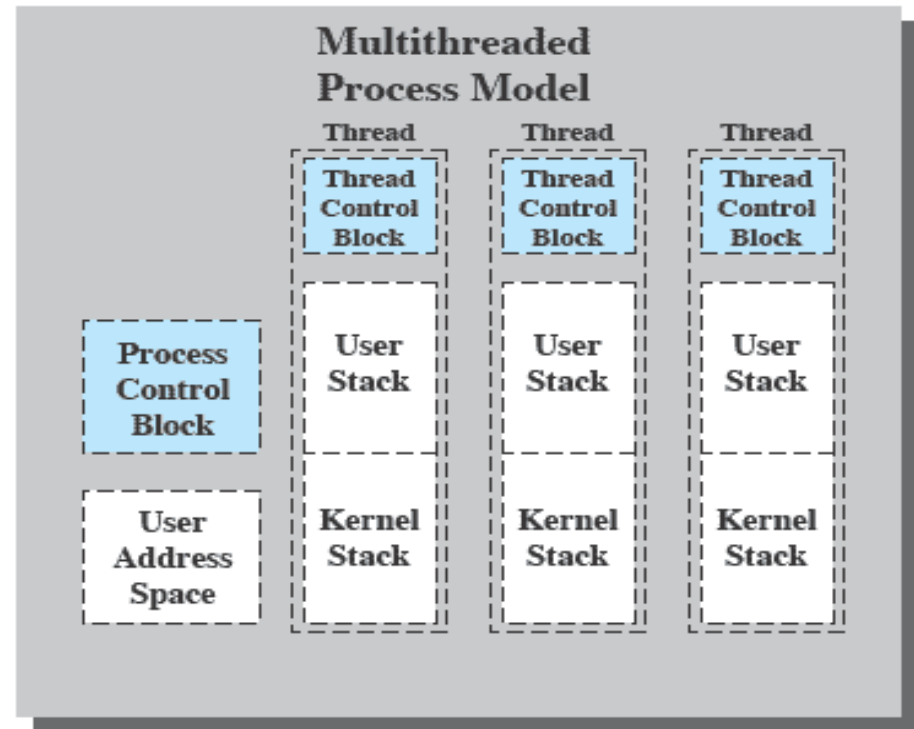
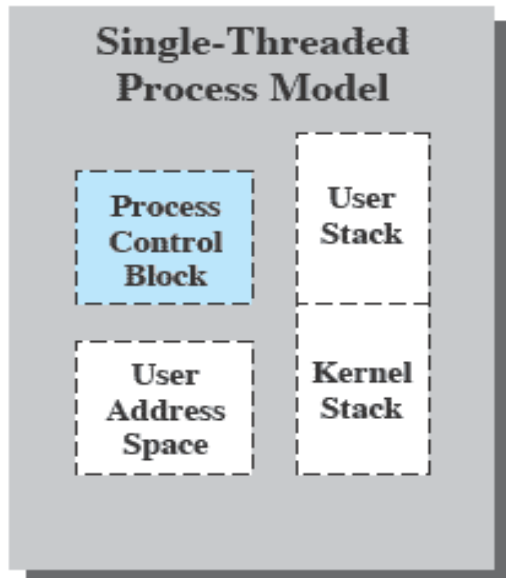
- ❑ Different part of a program may do different things and they can be executed **concurrently** to **improve response time** (or completion time).
 - Example: one thread may do a processor-bound task like rendering an image, while another thread responds to user interaction in the same program.
- ❑ If there is an interaction between different parts of the programs — **concurrency control** need to be applied.
- ❑ Example: accessing and modifying a common variable — **mutual exclusion** need to be satisfied.

Attributes of a Thread

Each thread has:

- an execution state (Running, Ready, etc.)
- saved thread context when not running
- an execution stack
- some per-thread static storage for local variables
- access to the memory and resources of its process (all threads of a process share this)

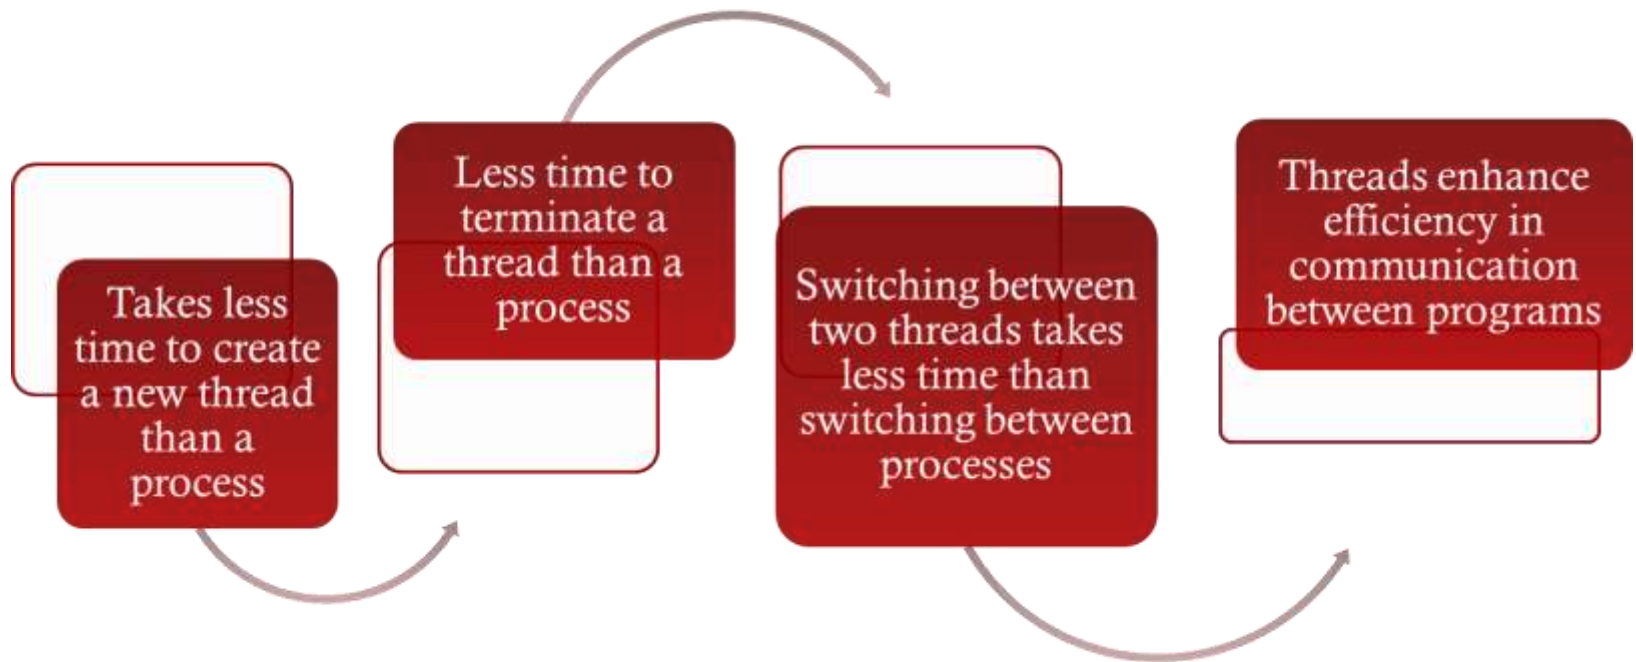
Threads vs. Processes



Single Threaded and Multithreaded Process Models

Benefits of Threads

Threads == 'lightweight processes'



More on Threads

- ❑ For an OS that supports threads, **scheduling** and **dispatching** is done on a thread basis.
- ❑ Most of the **state information** dealing with execution is maintained in **thread-level** data structures:
 - Suspending a process involves suspending all threads of a process.
 - Termination of a process terminates all threads within the process.

All threads within a process share the same address space.

What are the *thread states*?

Thread Execution States

Thread States

- ❑ The key states for a thread:
 - RUNNING
 - READY
 - BLOCKED

New threads

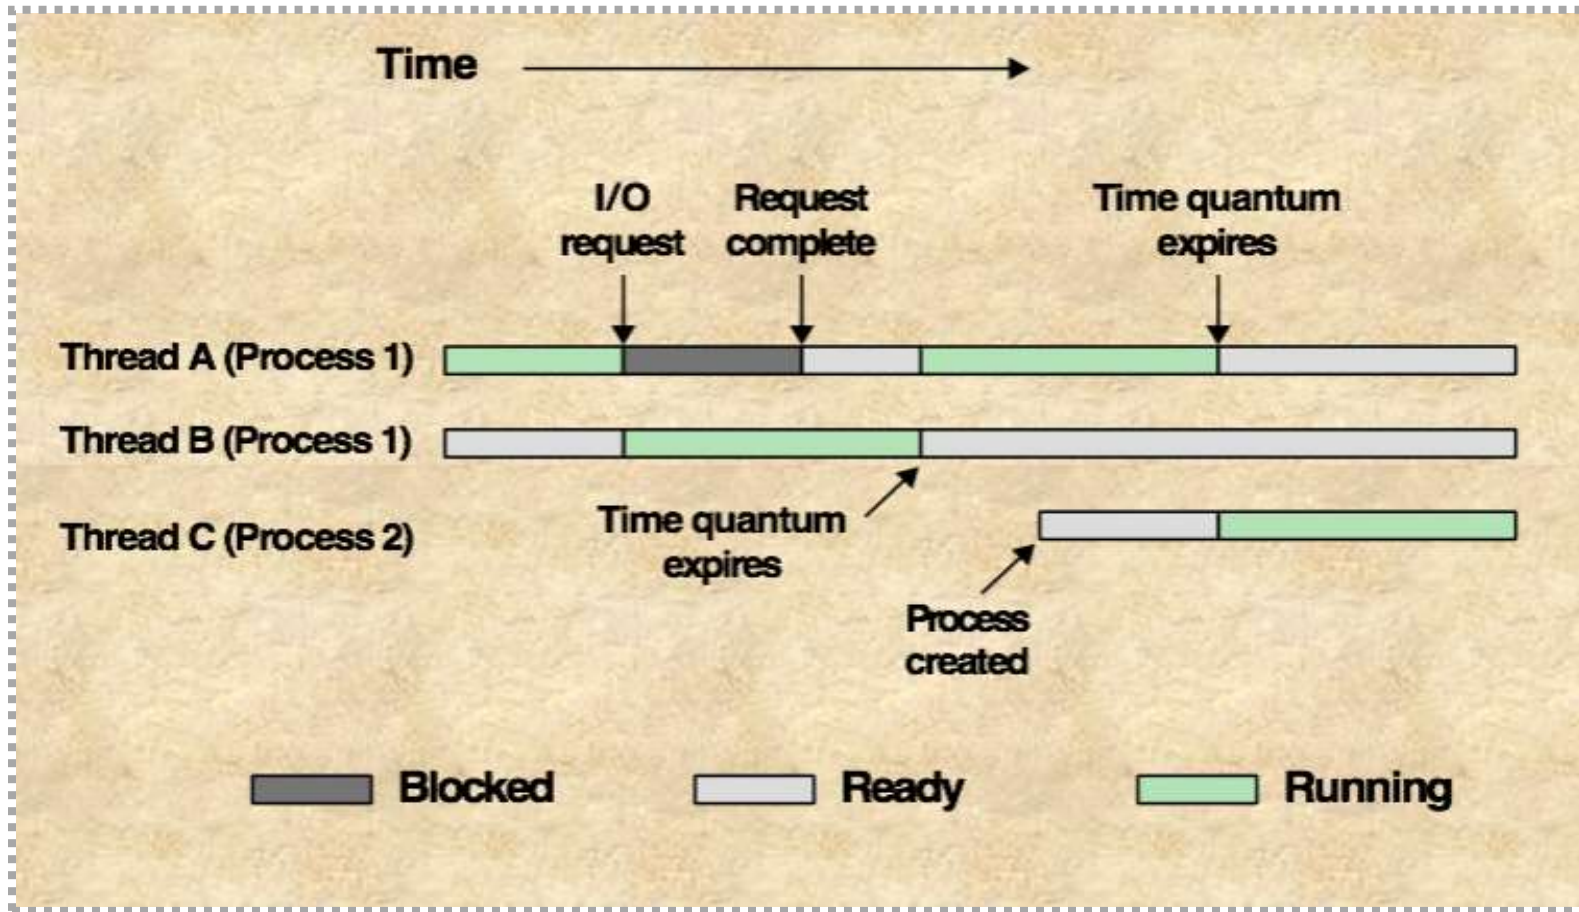
A BLOCKED event occurs — move the thread to the READY queue

State Transition

- ❑ Thread operations associated with a change in thread state:
 - Spawn
 - Block
 - Unblock
 - Finish

Wait for an event

Multi-threading on a Uniprocessor



e.g. One thread may run while another thread is blocked.


Thread Synchronisation

- ❑ It is necessary to **synchronise** the activities of various threads.
- ❑ All threads of a process share the same address space (e.g. global variables) and other system resources.
- ❑ Any alteration of a resource by one thread affects the other threads in the same process.

User-Level Thread (UTL) and Kernel-Level Thread (KLT)

User-Level Threads (ULTs)

- ❑ All thread management is done by the application.
- ❑ The kernel is not aware of the existence of threads.
- ❑ Any application can be programmed to be multi-threaded by using a **threads library**.
 - Even if OS does not support threads.



Pure user-level

ULTS: Advantages

Thread switching does not require kernel mode privileges

Scheduling can be application specific


ULTs can run on any Operating System

ULTS: Disadvantages

- ❑ In a typical OS, many system calls are *blocking*.
- ❑ When a ULT executes a system call, not only is that thread gets blocked, but all of the threads within the process are also blocked.
- ❑ In a pure ULT strategy, a multi-threaded application cannot take the full advantage of multiprocessing.

Kernel-Level Threads (ULTs)

- ❑ Thread management is done by the kernel.
- ❑ No thread management is done by the application — through API to the kernel thread facility.
- ❑ Example: Windows, Linux



Pure kernel-level

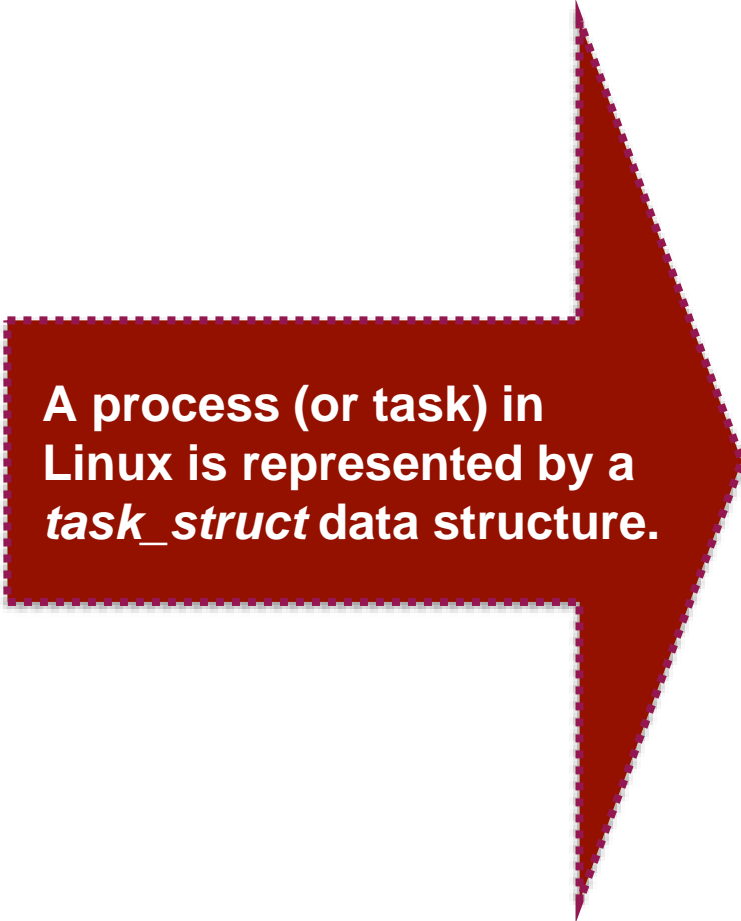
KLTS: Advantages

- ❑ The kernel can simultaneously schedule multiple threads from the same process on **multiple processors**.
- ❑ If one thread in a process is blocked, the kernel can schedule another thread of the same process.
- ❑ The kernel routines can also be multi-threaded.

KLTS: Disadvantages

- ❑ The transfer of control from one thread to another thread within the same process requires a **mode switch** to the kernel.
 - Some overhead here.

How do Unix/Linux systems manage threads?



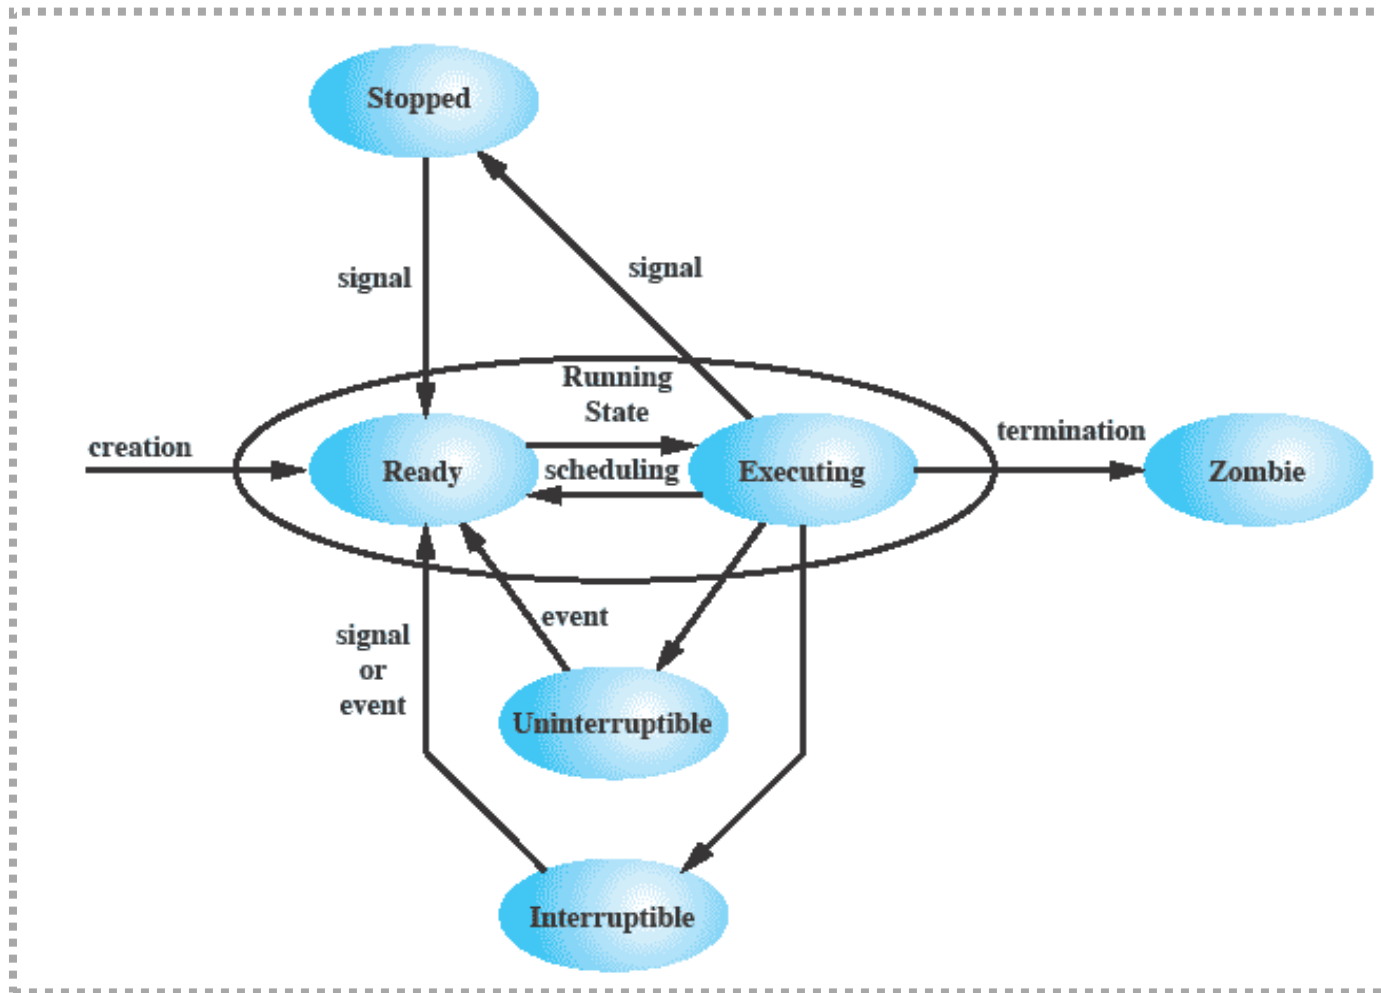
A process (or task) in Linux is represented by a *task_struct* data structure.

- Execution stage
- Scheduling information
- Identifiers (IDs)
- Links to parents/children/siblings
- Timers
- File system
- Address space
- Processor-specific context information

PTHREADS

- ❑ Threads in Linux are known as **pthread**s. Managed through a separate API.
- ❑ The **pthread** library must be included and linked into the program in order to use threads.
 - **#include <pthread.h>**
 - Add **-lpthread** to the end of the **gcc** command to **link** the program against the pthread library
- ❑ **pthread_create()** – spawn a new thread
- ❑ **pthread_join()** – wait for another thread to terminate

Linux: Process/Thread Model



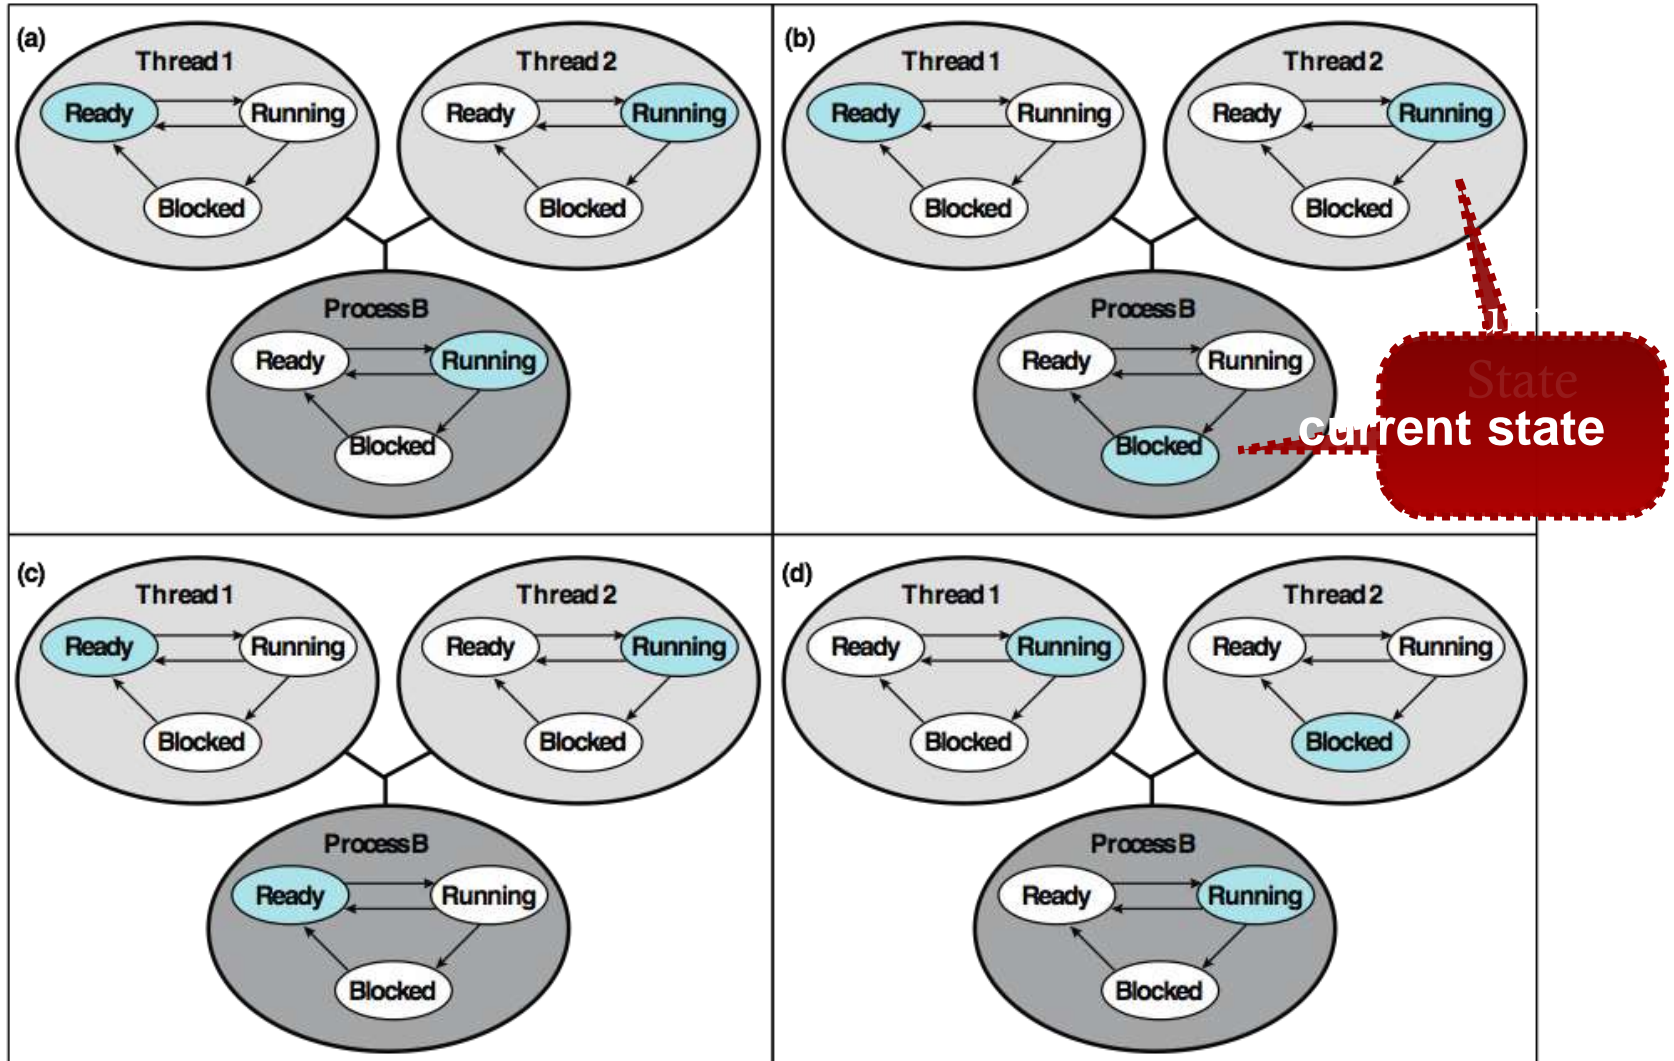
Summary of Lecture 4 (Part A)

- ❑ The concept of **process** is related to resource ownership.
- ❑ The concept of **thread** is related to program execution.
- ❑ In **multi-threaded** system, multiple concurrent threads may be defined with a single process.
- ❑ Two types of threads: **user-level** and **kernel level**.

Reading from Stallings, Chapter 4: 4.1, 4.2 and 4.6


SUPPLEMENTARY SLIDES

Thread Scheduling and Process Scheduling




Combined Approaches (ULT and KLT)

- ❑ Thread creation is done in the user space.
- ❑ Bulk of **scheduling** and **synchronisation** of threads are by the application.
- ❑ Multiple ULTs from a single application is mapped onto smaller (or equal) number of KLTs.
- ❑ Example: Solaris (Unix)



Linux: Namespaces

- ❑ A **namespace** enables a process to have a different view of the system than other processes that have other associated namespaces.
- ❑ One of the overall goals to support the implementation of **control groups** (*cgroups*).  **virtual machine**
- ❑ A tool for lightweight *virtualisation* — provides a process or group of processes with the illusion that they are the only processes on the system.
- ❑ Six namespaces in Linux:

mnt

pid

net

ipc

uts

user

Relationships between Threads and Processes

Threads:Processes	Description	Example Systems
1:1	Each thread of execution is a unique process with its own address space and resources.	Traditional UNIX implementations
M:1	A process defines an address space and dynamic resource ownership. Multiple threads may be created and executed within that process.	Windows NT, Solaris, Linux, OS/2, OS/390, MACH
1:M	A thread may migrate from one process environment to another. This allows a thread to be easily moved among distinct systems.	Ra (Clouds), Emerald
M:N	Combines attributes of M:1 and 1:M cases.	TRIX

Linux: Threads

