

VIRTUAL MEMORY

LECTURE 9 / FIT2100 / SEMESTER 2 2019

WEEK 10



VIRTUAL MEMORY

INTRODUCTION


LEARNING OUTCOMES

- Understand why **virtualised** memory is useful
- Discuss the differences between paging and segmentation
- Understand how to translate a logical memory address into a physical address
- Describe the most common page-replacement algorithms.

READING

- Stallings: Chapter 7, Chapter 8.

WHAT IS VIRTUAL MEMORY?



LET'S RECAP...

WHAT IS PHYSICAL MEMORY?

WHAT IS MAIN MEMORY?



A LONG, LINEAR LIST OF ADDRESSABLE BYTES.

- A physical memory address corresponds **directly** to a byte storage location in hardware.

VIRTUAL MEMORY IS...

AN ABSTRACTION OF PHYSICAL MEMORY

LOGICAL ADDRESSING

- Instead of dealing with physical memory addresses directly, the OS gives programs **logical addresses** to work with.
- Logical memory addresses are then translated into **physical addresses**.

THE 'REAL' MEMORY ADDRESS IS HIDDEN FROM THE PROGRAM.

- The logical address is also known as a **virtual address**.
- All memory address pointers you've worked with in C have been virtual addresses.

BUT WHY?

WHY VIRTUALISE THE ADDRESS SPACE? (1/3)

EFFICIENCY, SECURITY, CONCURRENCY, FLEXIBILITY, ...

- Efficient allocation of resources:
 - Main memory is a limited, expensive resource.
 - Buy more memory? But most of it will be un-used most of the time...
 - Why not swap parts of memory out to secondary storage when not required?
 - Divide memory into either **pages** or **segments** that can be swapped in and out... (later in lecture)
- Make better use of available memory:
- More processes can be run at a time
- A process can be run even if larger than **all** of main memory.

BUT WHY?

WHY VIRTUALISE THE ADDRESS SPACE? (2/3)

EFFICIENCY, SECURITY, CONCURRENCY, FLEXIBILITY, ...

- Enables protection against unauthorised access to another process's memory space
- Enables **shared memory** among related processes
 - A segment of memory can be 'mapped' into the memory space of multiple related processes
 - e.g. multiple processes can share certain variables
 - Multiple instances of the same program can share code.
- e.g. When **forking** or running multiple instances of the **same program**, the program code only needs to be loaded once.
- Each process sees its own **logical** copy of the program code, but the logical segments of memory are translated to a common physical location.
- This is why you can't modify hard-coded string literals in C. They are part of the program code itself, which may be mapped into multiple instances' memory spaces.

```
char* str = "Hi!";  
str[0] = 'h';  
//segmentation fault!
```

BUT WHY?

WHY VIRTUALISE THE ADDRESS SPACE? (3/3)

EFFICIENCY, SECURITY, CONCURRENCY, FLEXIBILITY, ...

- Relocation:
 - If a program must be swapped out to secondary storage, and is swapped back into main memory later on...
 - The same location in main memory might not be available.
 - Might be relocated to a different part of physical memory.
 - Relocating a program means the physical memory addresses are all different.
- All the program's instructions and variables now have different physical memory addresses.
- Dealing with logical addresses rather than physical addresses means the program's pointers remain valid
- The program can continue to run normally regardless of where it is located in physical memory.
 - Program does not depend on being loaded into a particular location.

SWAPPING


MAIN MEMORY ↔ SECONDARY STORAGE

WHAT IS SWAPPING?

- Moving chunks of data between main memory and secondary (disk) storage.
- Process data which is not currently needed can be moved out to the disk
- Moved back into physical memory when needed

IS IT USEFUL?

- Allows more complete utilisation of main memory
- But secondary storage is very **S L O W**
- Need to avoid swapping too often.
- **Thrashing**: a condition where the time spent of swapping data exceeds time spent executing instructions.



TWO APPROACHES

TO VIRTUAL MEMORY ALLOCATION


PAGING

Memory is **partitioned** into **fixed sized** chunks

SEGMENTATION

Memory is managed as multiple **segments** of **different sizes**

PAGED VIRTUAL MEMORY



```
01010000 01101100  
01100101 01100001  
01110011 01100101  
00100000 01100111  
01101001 01110110  
01100101 00100000  
01100111 01101111  
01101111 01100100  
00100000 01010011  
01000101 01010100  
01010101
```

FRAMES AND PAGES

PAGED VIRTUAL MEMORY

PAGES FIT INTO FRAMES

- **Physical memory** is partitioned into equal sized **frames**.
- A **page** is a chunk of data that can fit into a frame
- All frames are of equal size, and a page is the same size as a frame.
- Some pages might be in physical memory (loaded into frames)
- Other pages might be **swapped out** to secondary storage (taken out of frames)
 - To make way for other pages to be swapped in.
- The OS kernel maintains a **page table** to keep track of which pages are in which frames.



FRAGMENTATION

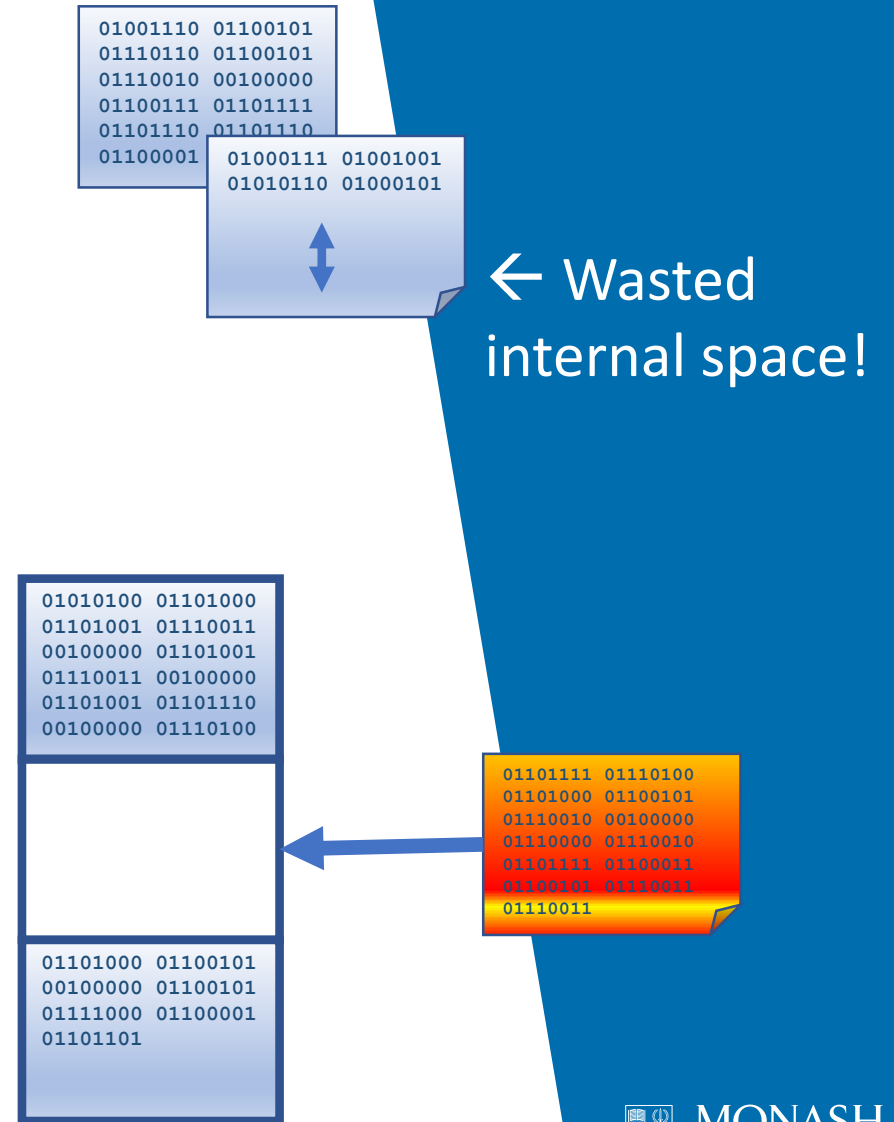
PAGED VIRTUAL MEMORY

INTERNAL FRAGMENTATION

- A process may occupy one or more frames in memory.
- If an entire page is not needed, the space cannot be used by another process
- Since all pages are fixed size, any excess space is wasted.

NO EXTERNAL FRAGMENTATION

- Can space **between** page frames be wasted?
- No. All pages are equally sized, so any empty frame in memory can **always** be allocated a whole page.
 - ➔ no external fragmentation
- ...but this is not true for memory systems that use **segmentation** (see later)



LOGICAL ADDRESSES

PAGED VIRTUAL MEMORY

THE MEMORY ADDRESS CONTAINS A PAGE NUMBER AND AN OFFSET

Example:

000000000010100000000000000010000

Page number 10 Offset 16 bytes

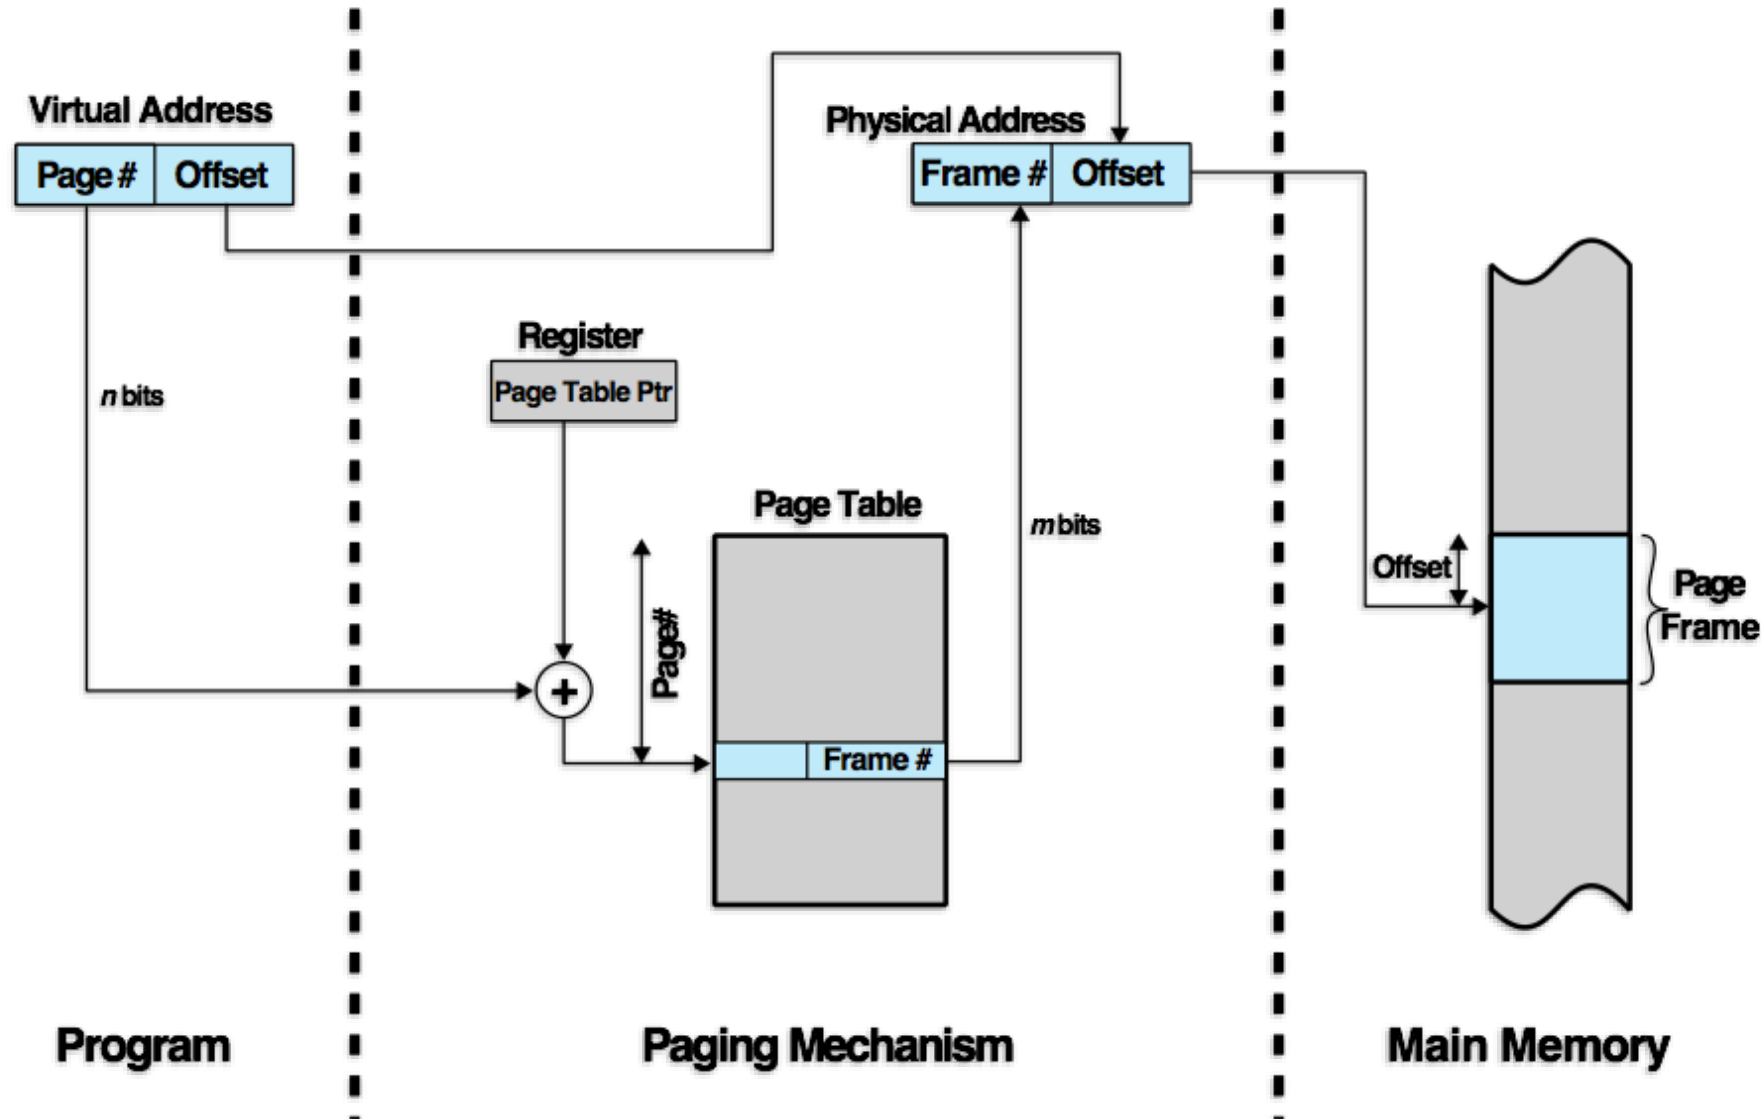
The most significant bits in the **logical address** are used for the page number. The remaining bits are an **offset** from the start of the page (i.e. position inside the page)

The address is structured so the **number of offset bits** matches the addressable range of the **page size**. Different systems may use a different page size.

- It is not possible to specify an offset that goes 'out of bounds' of the specified page.

ADDRESS TRANSLATION (PAGING)

Change (logical) page number to (physical) frame number → physical address.



WHAT IS A PAGE FAULT?

PAGED VIRTUAL MEMORY

WHAT HAPPENS IF A PAGE DOES NOT HAVE A FRAME IN THE PAGE TABLE?

- It means the page does not have a physical address in memory

‘PAGE FAULT’!

- The requested **logical** memory address is **valid**, but the page is in **secondary storage**. Must be loaded in to a frame.

IS A PAGE FAULT ‘BAD’?

- If all the frames are already full, another page needs to be swapped out of main memory to make room for the needed one.
- We need a ‘page replacement algorithm’ to determine which page should be swapped out to make way for the one that’s needed.
- **NOTE:** even if there **is** an empty frame to load the needed page into, it is still a **page fault** if the page was not found in physical memory.

PAGE REPLACEMENT ALGORITHMS

PAGED VIRTUAL MEMORY

PAGE FAULT == PAGE NOT FOUND IN PHYSICAL MEMORY

- If we get a page fault, we need to load a page from disk into a frame in RAM.
- What if all the frames are already full?
- Need a strategy for knowing which page to swap out.
- AIM:
 - Swap out a page that won't be needed again for a long time
 - PREVENT **THRASHING**
 - Try to make a 'fast guess'. Don't try to be 'too clever'.
 - MORE ACCURATE PREDICTION → MORE COMPUTATION TIME NEEDED
 - LESS ACCURATE PREDICITON → MORE PAGE SWAPS NEEDED
 - TRY TO FIND A BALANCE.

FIRST IN FIRST OUT (FIFO)

SIMPLEST REPLACEMENT POLICY TO IMPLEMENT

PAGE THAT HAS BEEN IN MEMORY THE LONGEST IS REPLACED

Treat frames allocated to processes as a 'queue' for replacement.


Pages are removed in round-robin style.

EXAMPLE

Page address
stream

2 3 2 1 5 2 4 5 3 2 5 2

FIFO



F = page fault

Total **9 page faults**.

(or 6 page faults after all frames filled)

LEAST RECENTLY USED (LRU)

DIFFICULT TO IMPLEMENT

REPLACE THE PAGE THAT HAS NOT BEEN REFERENCED FOR THE LONGEST TIME

By the **principle of locality**, pages that have been referenced recently are likely to be referenced again in near future.


A **lot of overhead** to maintain/check the time since last reference for every page.

EXAMPLE

Page address
stream

2 3 2 1 5 2 4 5 3 2 5 2

LRU



F = page fault

Total **7 page faults**.

(or 4 page faults after all frames filled)

OPTIMAL POLICY (OPT)

THEORETICAL BEST CASE

REPLACE THE PAGE THAT WILL NOT BE REFERENCED FOR THE LONGEST TIME

Produces the smallest possible number of page faults.

Impossible to implement in practice: the operating system cannot see the future.


- If we knew everything the process would do in the future, we would not need to run the program!

EXAMPLE

Page address
stream

2 3 2 1 5 2 4 5 3 2 5 2

OPT



F = page fault

Total **6 page faults**.

(or 3 page faults after all frames filled)

PAGE REPLACEMENT ALGORITHMS

SUMMARY


PERFORMANCE

	PAGE FAULTS	COMPLEXITY
FIFO	9	Simplest
LRU	7	Complex
OPT	6	'Impossible'

PAGE REPLACEMENT ALGORITHMS ARE NOT PERFECT

- Page replacement is a trade-off between **minimising** page faults and **minimising** computational complexity.
- Page replacement can happen **very often** (whenever data is swapped in). Identifying a page to replace must to be done in **as few instructions as possible** while still giving a 'good enough' result.
 - So that the program's instructions can be executed instead!

SEGMENTED VIRTUAL MEMORY




SEGMENTS

SEGMENTED VIRTUAL MEMORY

WHAT IS A SEGMENT?

- A process is allocated one or more logical **segments** of memory.
- A segment can be swapped in or out, just like a page.
- Segments can vary in size, up to a defined maximum size.
- There are no fixed 'frames' in physical memory: segments are placed where space is available.
- Swapping out a segment leaves a 'hole' where another segment may fit.



FRAGMENTATION


SEGMENTED VIRTUAL MEMORY

NO INTERNAL FRAGMENTATION

- A segment may grow or shrink dynamically depending on required size. There is no 'wasted' space within a segment that cannot be used.

EXTERNAL FRAGMENTATION

- Can space **between** segments be wasted? **YES.**
 - A small segment may be swapped out (or a small process may exit the system).
 - The **hole** left behind may be too small to reasonably fit another segment
 - Even though the space is available, it is not used.
- **Compaction** is the technique of periodically re-arranging segments to remove gaps. Computationally expensive.



PLACEMENT ALGORITHMS

SEGMENTED VIRTUAL MEMORY

Best-fit

- chooses the block that is **closest in size** to the request

First-fit

- begins to **scan memory from the beginning** and chooses the first available block that is large enough

Next-fit

- begins to **scan memory from the location of the last placement** and chooses the next available block that is large enough

LOGICAL ADDRESSES

SEGMENTED VIRTUAL MEMORY

THE MEMORY ADDRESS CONTAINS A SEGMENT NUMBER AND AN OFFSET (Just like paging)

Example:

000000001100000000000000000010001

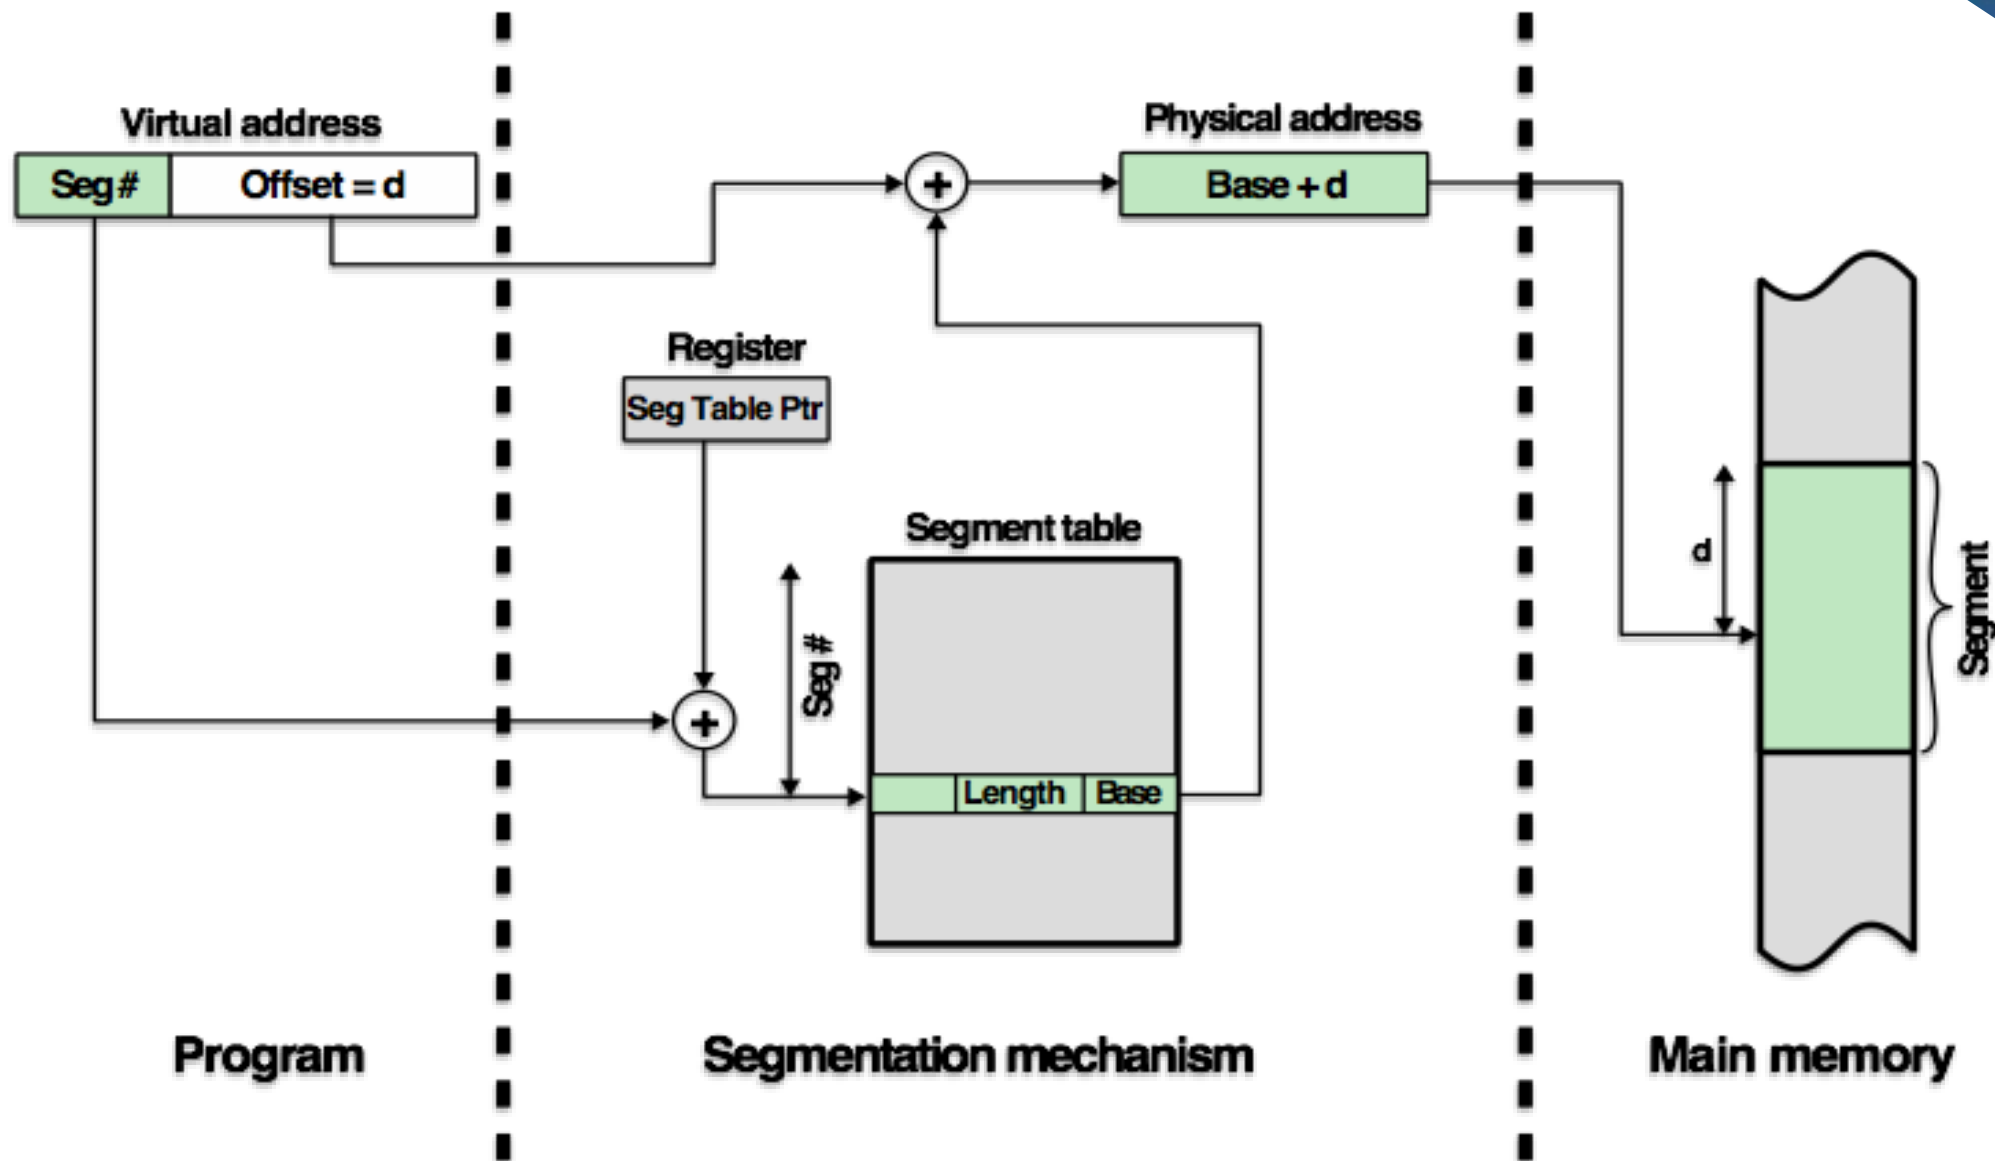
Seg number 12 Offset 17 bytes

The most significant bits in the **logical address** are used for the **segment number**. The remaining bits are an **offset** from the **start** of the segment in physical memory.

Segments can be different sizes. **If the offset goes past the end of a segment**, the logical address is invalid.

‘SEGMENTATION FAULT’!!! ← Process is usually terminated.

ADDRESS TRANSLATION (SEGMENTATION)



SHARED MEMORY

A SEGMENT IN MEMORY THAT IS ALLOCATED TO MULTIPLE PROCESSES

- A memory segment may be shared among multiple processes
- Like with multi-threading, multiple processes can access each other's memory
- Unlike with multi-threading, it does not happen automatically.
 - Must manually request a new block of shared memory to be allocated.
 - Required number of bytes must be specified.
 - Must manually assign a data structure to the shared block of bytes.
- There is no concurrency protection offered. Mutual exclusion mechanisms like semaphores must be used.
- Shared memory is an important tool for inter-process communication. Even unrelated processes can be written to share memory.

SUMMARY (LECTURE 9)

VIRTUAL MEMORY

- We now understand why **virtualised** memory is useful
- Paging and segmentation are different approaches to virtual memory and create different challenges.
- We understand how to translate a logical memory address into a physical address by adding the offset to the starting address of the page/segment in physical memory.
- We have discussed three different page-replacement algorithms.
- **Next week: Inter-process communication**