



FIT2100 Practical #2

Managing Larger C Programs and Running I/O System Calls in C

Week 4 Semester 2 2019

August 13, 2019

Revision Status:

\$Id: FIT2100-Practical-02.tex, Version 2.0 2018/08/07 17:30 Jojo \$

Written by Jojo Wong, updated by Daniel Kos, Aug 2019.

Acknowledgements

The content presented in Sections 2.1 and 2.2 in this practical were adapted from the courseware of FIT3042 prepared by Robyn McNamara.

The content presented in Section 2.3 and the practical tasks (Section 3) were adapted from David Curry's texts.

- David A. Curry (1989). *C on the UNIX System*, O'Reilly.
- David A. Curry (1996). *UNIX Systems Programming for SVR4*, O'Reilly.

Contents

1	Background	4
2	Pre-lab Preparation (3 marks)	4
2.1	Managing C programs	4
2.1.1	Organising your source files	4
2.1.2	Building complex C programs	5
2.1.3	The <code>makefile</code> format	6
2.1.4	Running the <code>make</code> command	7
2.1.5	Defining macros in <code>makefile</code>	7
2.1.6	Pre-class exercise (3 marks)	8
2.2	I/O System Calls in C	9
2.2.1	File descriptors	9
2.2.2	Opening and creating files	10
2.2.3	Reading and writing files	12

2.2.4	Using write() with strings	12
2.2.5	Moving around files	13
3	Practical Tasks (7 marks)	15
3.1	Task 1 (2 mark)	15
3.2	Task 2 (2 mark)	17
3.3	Task 3 (3 marks)	18
3.4	Task 4 (0 mark)	18
4	Postscript: redirects and pipes	18

1 Background

This practical has two objectives. Firstly, you will learn how to manage complex C programs. Secondly, you will learn how to perform low-level I/O systems calls in C.

There are some pre-lab preparation (Section 2) that you should complete before attending the lab. The practical tasks specified in Section 3 are to be assessed in the lab.

Before you attempt any of the tasks in this prac, create a folder named PRAC02 under the FIT2100 folder (`~/Document/FIT2100`). Save all your source files under this PRAC02 folder.

2 Pre-lab Preparation (3 marks)

2.1 Managing C programs

Often times, working with a small programming project, you could probably get away with putting all your code in one source file. For much larger projects, you would need to organise your functions together into various source files by grouping related functions together.

A general *rule of thumb* is that your functions should not be much larger than a terminal screenful (approximately 25 lines of code), excluding header comments. If your functions become larger than this, you should consider breaking them down into smaller functions. (Your function names should be descriptive to help other coders to navigate your source files.)

2.1.1 Organising your source files

One characteristic of the C Programming Language is that it encourages the idea of 'separate compilation.'

Instead of putting all the code in one file and compiling that one file, C allows you to write many `.c` files and compile them separately when necessary. The usual way to group related C source files is to put them under the same directory.

What goes into the `.c` file? Anything that causes the compiler to generate code should go into the `.c` source file. A `.c` source file usually consists of:

- implementation of functions (i.e. function bodies)

- declaration of global variables — in order to set aside memory for global variables

Anything that is intended as a message to the compiler should go into the `.h` header file, which include:

- function prototypes
- struct definitions
- typedef statements
- #define statements

Note: Only `.h` files should ever be included using the `#include` statement — you should never include `.c` files.

For example, if your program has a `.c` file named `arithmetic.c`, you would place the function prototypes and other related statements in a header file named `arithmetic.h`, and `#include "arithmetic.h"` both at the top of `arithmetic.c` and at the top of any other file that needs access to those function prototypes. Including just the header information in each file rather than the entire functions themselves avoids a problem where the C compiler 'sees' the same function redefined multiple times.

2.1.2 Building complex C programs

A C program can potentially consist of hundreds of source files (both `.c` and `.h` files), and possibly spanning across many directories.

The standard command given below is used to compile a complex C program with multiple source files. However, *this approach does not scale!*

```
1 $ gcc sourcefile1.c sourcefile2.c ... -o myprogram
```

Whenever the program is rebuilt, this command recompiles every source file, which can be time consuming. It should only be recompiling those source files that have been modified instead of every single file. The solution to this is by utilising the `make` command in C.

`make` is a utility tool in C that automates the building of program binaries from source files. A number of built-in rules are defined, such that `make` knows how to compile source code (`.c` files) into object code (`.o` files). `make` requires a special file called `makefile` that specifies the source files and the targets (which include the executable code) to be build.

2.1.3 The makefile format

A makefile typically consists of one or more entries. Each entry consists of the following:

- a target (which usually is a file but not always);
- a list of dependencies (files which the target depends on);
- and the commands to execute based on the target and its dependencies.

The basic format for each entry in a makefile looks like below:

```
1 <target>: [ <dependency> ]*  
2 [ <TAB> <command> <ENDLINE> ]+
```

(**Note:** There must be a <TAB> character at the start of each command defined under the target. Indentation using spaces will *not* work in a makefile.)

An example of a makefile:

```
1 sourcefile1.o: sourcefile1.c headerfile1.h  
2 gcc -Wall -c sourcefile1.c
```

In the example above, the `-c` option is needed to create the corresponding object file (`.o`) for the given `.c` file. The `-Wall` option (stands for *Warnings: all*), which you might not have seen before, is a way to turn on extra warnings in `gcc`, which can be very useful for finding hidden problems in your code.

How and when is a target constructed? Each entry in the makefile defines *how* and *when* to construct a target based on its dependencies. The dependencies are used by the makefile to determine when the target needs to be reconstructed (by re-compiling the source code).

Each file has a *timestamp* that indicates when the file was last modified.

The `make` command first checks the timestamp of the file (i.e. the target) and then the timestamp of its dependencies (which are also files). If the dependent files have a more recent timestamp than the target file, the command lines defined under the entry of that target are executed to update the target.

Note: The updates are done recursively. If any of the dependencies themselves are also targets (i.e. there is a separate entry in the makefile for each of the dependencies), the `make`

command will check whether these dependent files need to be updated, before updating the initial target.

How to create the executable code using targets? The purpose of having the makefile is to create the executable code (i.e. a.out). The executable file is often the first target.

```
1 myprogram: sourcefile1.o sourcefile2.o sourcefile3.o
2 gcc -Wall sourcefile1.o sourcefile2.o sourcefile3.o -o myprogram
```

2.1.4 Running the make command

Once you have a makefile written, you can make use of the `make` command to update your files whenever you make changes. There are two ways to run `make` as shown below.

```
1 $ make
2 $ make -f <makefilename>
```

If you just type `make` on the command prompt, the file called `makefile` in the current directory will be interpreted and the commands of the first target will be executed, provided that the first target has dependencies listed.

Alternatively, if you have more than one `makefile` with different file names, you will run the `make` command with the `-f` option by supplying the file name of the specific `makefile` that you would like to run.

2.1.5 Defining macros in makefile

You can use *macros* in your `makefile`, which allow you to define variables that can be substituted.

Macros have a similar syntax to shell variables found in `bash`, where you can use the `"=`" operator to set a value to the macro, and use the `"$"` operator when using the value of the macro. (Note that the parentheses are required if the macro name has more than one character.)

Important: Don't get confused with the shell command prompt which can also be represented by a `"$"` symbol.

```

1 CC = gcc
2 CFLAGS = -Wall -c
3
4 sourcefile1.o: sourcefile1.c headerfile1.h
5 $(CC) $(CFLAGS) sourcefile1.c

```

The example above demonstrates two common macros CC and CFLAGS. Below are some of the other common examples of macros.

Macro	Description
CC	the name of the compiler (gcc)
DEBUG	the debugging flag (-g)
CFLAGS	the flags used for compiling (e.g. -c, -Wall)
LFLAGS	the flags used for linking (e.g. -L, -I)
\$@	the file name of the target
\$^	the list of dependencies
\$<	the first item in the dependencies list

Putting it all together Here is a sample of makefile for creating the executable program called myprogram.

```

1 OBJS = sourcefile1.o sourcefile2.o sourcefile3.o
2 CC = gcc
3 CFLAGS = -Wall -c
4 LFLAGS = -Wall
5
6 myprogram: $(OBJS)
7     $(CC) $(LFLAGS) $(OBJS) -o myprogram
8
9 sourcefile1.o: sourcefile1.c headerfile1.h
10    $(CC) $(CFLAGS) sourcefile1.c
11
12 sourcefile2.o: sourcefile2.c headerfile2.h
13    $(CC) $(CFLAGS) sourcefile2.c
14
15 sourcefile3.o: sourcefile3.c headerfile3.h
16    $(CC) $(CFLAGS) sourcefile3.c

```

2.1.6 Pre-class exercise (3 marks)

Refer to Task 3 that you have completed for the Week 1 Tutorial. First, organise your code into three separate files: main.c, arithmetic.c, and arithmetic.h.

- `main.c`: consists of the `main()` function
- `arithmetic.c`: consists of the implementation of the four arithmetic functions (i.e. function bodies)
- `arithmetic.h`: consists of the function prototypes of the four arithmetic functions

Then, create a `makefile` and try to compile using the `make` command.

2.2 I/O System Calls in C

Note: You should complete the reading of this section before attempting the practical tasks in the next section.

The Standard I/O Library (`stdio`) provides a collection of *high-level* routines to perform input and output (I/O). This enables C programmers to perform reading and writing data easily and efficiently.

In general, the `stdio` routines perform three important functions for the programmers:

- Buffering on input and output data is performed automatically;
- Input and output conversions are performed using routines like `printf` and `scanf`;
- Input and output are also automatically formatted.

However, these functions such as buffering and I/O conversions are not always applicable. In the event where performing I/O directly from a device such as a disk drive, programmers would need to be able to determine the actual buffer size to be used instead of relying on the `stdio` routines.

Thus, a direct interface to the operating system is desirable and this can be achieved by issuing *system calls* or using the *low-level* I/O interface.

2.2.1 File descriptors

When dealing with the low-level I/O interface, each of low-level I/O routines requires a valid “file descriptor” to be passed to them — which is used to reference a file to an open stream for I/O.

A file descriptor is simply a small integer, which can be allocated from a file index table maintained for each process (or program) by the operating system.

There are three pre-defined file descriptors. You do not need to open these yourself; the operating system provides these automatically:

- File descriptor 0: refers to standard input (usually the terminal keyboard)
- File descriptor 1: refers to standard output (usually the terminal screen)
- File descriptor 2: refers to standard error output (usually the terminal screen)

When a process (or program) begins with its execution, it starts out with three opened files — i.e. the three pre-defined file descriptors 0, 1, and 2. When the process opens its first file, it will then be attached with file descriptor 3.

2.2.2 Opening and creating files

In order to read from or write to a file (including the standard input and output), that file must first be opened. The routine (or function) used to *open* a file or to *create* a file for reading and/or writing is called `open()`.

The `open()` function takes three arguments:

- a character string consists of the path name of a file to be opened;
- a set of flags or integer constants specifying how the file is to be opened;
- an integer *mode* used to specifying the access permission when creating a new file.

The man page for `open` (`man 2 open`) illustrates how this function is declared internally inside the system libraries, and what headers you must `#include` to use it.

```
1 #include <sys/types.h>
2 #include <sys/stat.h>
3 #include <fcntl.h>
4
5 int open(const char *pathname, int flags);
```

The second argument is constructed by OR-ing together¹ a number of flags or constants (see below) defined in the header file `sys/fcntl.h` for System V systems, and `sys/file.h` for BSD (Berkeley Software Distribution) systems.

A subset of the flags (constants) that control how a file is to be opened are shown in the table on the next page. Note that the first three flags are mutually exclusive.

The `open` routine returns a file descriptor for the file if it is opened successfully. Otherwise, the value of `-1` is returned if it fails to open that file. An error code indicating the reason for failure is stored in the external variable `errno` defined in the header file `errno.h`. You can print out the error message with the `perror` function.

Flag (Constant)	Description
<code>O_RDONLY</code>	The file is opened for read only
<code>O_WRONLY</code>	The file is opened for write only
<code>O_RDWR</code>	The file is opened for both read and write
<code>O_APPEND</code>	The file is opened in append mode
<code>O_CREAT</code>	Create the file if it does not exist; the mode argument must be supplied to specify the access permission on the file
<code>O_EXCL</code>	Check if the file to be created is already exists; if already exists, <code>open()</code> will fail
<code>O_NONBLOCK</code> or <code>O_NDELAY</code>	The file is opened in non-blocking mode
<code>O_TRUNC</code>	If the file is opened for writing, truncate the length of the file to zero

Often times, the opened file should be closed once a process (or program) has finished using that file. The routine to *close* a file is called `close()`, which takes in only one argument representing the file descriptor of the file to be closed. If the file was closed successfully, the value of `0` is returned; otherwise, the value of `-1` is returned if an error occurs.

```
1 #include <unistd.h>
2
3 int close(int fd);
```

¹Use the bitwise-or operator: `|` to combine multiple flag values together if needed.

2.2.3 Reading and writing files

With the low-level I/O interface, the `read()` and `write()` routines are used for reading from and writing to a file respectively.

```
1  #include <unistd.h>
2
3  ssize_t read(int fd, void *buf, size_t count);
4  ssize_t write(int fd, const void *buf, size_t count);
```

The `read()` attempts to read up to `count` bytes from the file referred to by the file descriptor `fd` into the buffer pointed by `buf`. The function returns the number of bytes actually read if no error occurs; otherwise the value of `-1` is returned and the external variable `errno` is set with the error code. If the return value is `0`, that indicates the end of the file has been reached and there is no data left to be read.

The `write()` writes up to `count` bytes from the buffer pointed by `buf` to the file referred to by the file descriptor `fd`. If the write is successful, the number of bytes actually written is returned; otherwise the value of `-1` is returned if an error occurs and `errno` is set accordingly. (If nothing was written, the return value is `0`.)

When using these low-level functions, the programmers have to decide on the size of the buffer in bytes (as indicated by the third argument `count` in both functions). Note that each time the `read()` or `write()` function is executed, the operating system will access the disk (or other I/O device). If the buffer size is set to `1`, meaning that you would only be able to read or write one byte (one character) at a time — the overall execution of your program is less efficient (especially when dealing with a large file) compared to reading or writing a whole block of characters at a time. In many cases the difference is not noticeable however.

2.2.4 Using `write()` with strings

While many C library functions use the C language convention of 'null-terminated' strings (using a `'\0'` character to mark the length of the string in memory), the underlying operating system does not assign any special meaning to the null `'\0'` character. Therefore, when making a `write()` system call to the operating system, you must explicitly specify the number of characters to be written, regardless of whether the string is already null-terminated.² The `strlen()` function defined in `<string.h>` can be useful here: this function counts the number of characters in a (null-terminated) string by finding the position of the null character.

²In some cases, such as image and video processing, the null character may even represent useful data rather than the end of the string

Common mistake: Beware of trying to use the `sizeof()` function to get the length of a string: depending on how your string is defined, you might be mistakenly getting the size of a `char*` (memory addresses are always 4 bytes in size on a 32-bit machine) or the size of an entire array, rather than the number of meaningful characters defined within it. In general, the `sizeof()` function is not the right tool for this job. It's only really meant for getting the memory size of data types rather than values themselves.

The following example shows how to use `write()` to print a string to standard output (which is file descriptor 1):

```
1 //requires inclusion of <string.h>
2
3 char* outputString = "Hello , world!\n";
4 write(1, outputString, strlen(outputString));
5
6 //can we put the whole thing in one line like with a printf()?
7
8 write(1, "This also works\n", 16);
9 /*...but note that hard-coding the length like this is
10 not considered ideal coding practice for readability! */
```

2.2.5 Moving around files

Often times, we need to be able to move around within a file to access the data stored at a specific location in the file. Each file is associated with a value, known as the *file offset*. It is often set to 0 indicating the beginning of the file when a file is opened or created.

The value of the file offset can be obtained or re-set by using the low-level routine `lseek()`.

```
1 #include <sys/types.h>
2 #include <unistd.h>
3
4 off_t lseek(int fd, off_t offset, int whence);
```

The `lseek()` function re-positions the file offset (`offset`) of a file referred to by the file descriptor `fd`, based on the position in the file specified by the last argument in the function `whence`. The values that can be assigned to `whence` are as follows:

Offset (Constant)	Description
SEEK_SET	The file offset is set to offset bytes from the beginning of the file
SEEK_CUR	The file offset is set to offset bytes from its current position
SEEK_END	The file offset is set to offset bytes from the end of the file

So, if you were to move to the beginning or the end of a file:

```
1  lseek(fd, 0, SEEK_SET);    /* move to the beginning of the file */
2  lseek(fd, 0, SEEK_END);    /* move to the end of the file */
```

To obtain the value of the current offset of the file:

```
1  off_t current;
2  current = lseek(fd, 0, SEEK_CUR);
```

Note that the `lseek()` returns the new offset value in bytes from the beginning of the file if it runs successfully. If there is an error, the value of `-1` is returned. (The `errno` variable is set to indicate the error.)

3 Practical Tasks (7 marks)

3.1 Task 1 (2 mark)

```
1  #include <sys/types.h>
2  #include <sys/stat.h>
3  #include <sys/file.h>    /* change to <sys/fcntl.h> for System V */
4  #include <unistd.h>
5
6  /*
7   * appendfile.c: append the contents of the first file to the second file
8   */
9  int main (int argc, char *argv[])
10 {
11     int n, infile, outfile;
12     char buffer[1024];
13
14     if (argc != 3) { //Q: what does argc (argument count) mean?
15         write(2, "Usage: ./appendfile file1 file2\n", 32);
16         exit(1);
17     }
18
19     /*
20      * Open the first file (file1) for reading
21      */
22     if ((infile = open(argv[1], O_RDONLY)) < 0) {
23         perror(argv[1]);
24         exit(1);
25     }
26
27     /*
28      * Open the second file (file2) for writing
29      */
30     if ((outfile = open(argv[2], O_WRONLY | O_APPEND)) < 0) {
31         perror(argv[2]);
32         exit(1);
33     }
34
35     /*
36      * CODE HERE: Copy data from the first file to the second file
37      */
38
39     /*
40      * Close the two files before exiting
41      */
42     close(infile);
43     close(outfile);
44     exit(0);
45 }
```

`appendfile.c` is a partially completed C program, that takes two file names — `file1` and `file2` — as the command-line arguments. It opens the first file (`file1`) for reading and the second file (`file2`) for writing. The contents of the first file (`file1`) is appended at the end of the second file (`file2`).

Your task is to complete the missing code segment in the program as indicated by the comment — `/* CODE HERE */`. To test the program, create two files and name them as `file1` and `file2`. The contents of each file is given below by using the `less` command in Unix.

```
1 $ less file1
2 Line 1: the first sentence in file1
3 Line 2: the second sentence in file1
4 Line 3: the third sentence in file1
5
6 $ less file2
7 Line 1: the first sentence in file2
8 Line 2: the second sentence in file2
9 Line 3: the third sentence in file2
```

Note: To run or test the program in Unix, use the following command:

```
1 $ appendfile file1 file2
```


3.2 Task 2 (2 mark)

Type out the given program below, and you may name it as `seekinfile.c`. Compile the program and run it. Your task is to describe what does the program do. Include your description with 2-3 sentences as a comment at the beginning of the program code.

```
1  #include <sys/types.h>
2  #include <sys/stat.h>
3  #include <sys/file.h>      /* change to <sys/fcntl.h> for System V */
4  #include <unistd.h>
5  #include <string.h>
6
7  struct record {
8      int  userid;
9      char username[6];
10 };
11
12 char *usernames[] = { "userA", "userB", "userC", "userD" };
13
14 int main(int argc, char *argv[])
15 {
16     int i, outfile;
17     struct record eachrec;
18
19     /*
20      * Open the file (recordfile) for writing
21      */
22     if ((outfile = open("recordfile", O_WRONLY | O_CREAT, 0664)) < 0) {
23         perror("recordfile");
24         exit(1);
25     }
26
27     for (i = 3; i >= 0; i--) {
28         /*
29          * Create a new record
30          */
31         eachrec.userid = i;
32         strcpy(eachrec.username, usernames[i]);
33
34         /*
35          * Write the record into the file
36          */
37         lseek(outfile, (long) i * sizeof(struct record), SEEK_SET);
38         write(outfile, &eachrec, sizeof(struct record));
39     }
40
41     close(outfile);
42     exit(0);
43 }
```

3.3 Task 3 (3 marks)

By modifying the program that you have understood in Task 2 (3.2), write a C program that reads the records from the `recordfile` in the following order: the second record (1), the last or fourth record (3), the first record (0), and the third record (2). Then print the records out on the terminal screen.

Don't forget that you have to first open the `recordfile` for reading; and use the low-level `read()` to read each of the records.

(Note: As long as your program is able to print out the username of each record, you have achieved the task given in this section.)

3.4 Task 4 (0 mark)

Try running NetData (found on your desktop under the Dev Tools directory) while your program is running. Locate the graph for disk I/O. Do you see any disk activity (in or out) when running your Task 2 program? Our files are quite small. How does it compare to the amount of disk I/O activity you see when loading a web browser or other large application?

4 Postscript: redirects and pipes

You might be wondering why Unix provides two different output streams (standard output, and also standard error). The answer is that Unix systems generally provide a lot of flexibility in how input and output are managed. The following operators are provided by the bash command shell, and work with any program, even if your program is not written to use them:

```
1 $ ./some_program > outputfile.txt
2 [Standard output produced by the program is redirected into a file instead]
3
4 $ ./some_program 2> errorfile.txt
5 [the program's error messages are redirected into a file instead of the
   terminal]
6
7 $ ./some_program < input_commands.txt
8 [the program takes its input from a file, not from the keyboard]
9
10 $ ./some_program | ./another_program
11 [the standard output from some_program is used as input for another_program
   . This is called a 'pipe'.]
```

These operators are invisible to the running program. The `bash` command shell does not pass them to your program as arguments.

Another example involves using the special `/dev/null` device. This is a special virtual device on a Linux system. Anything written to it is simply discarded. Suppose you want to hide error messages when running a certain program, and only want to see the meaningful output:

```
1 $ ./noisy_program 2> /dev/null
```